

Scaling our the Cambridge Investment Zone

AI-Native Infrastructure for Ecosystem Automation, Fintech Integration, and
Global Hub Connectivity

Dr. Yaniv Proselkov, Global Cambridge Institute

April 2026

Abstract

Cambridge produces world-class scientific research and startup activity, but still struggles to retain and scale the value created from that research. This paper is the third in a series by the Global Cambridge Institute: the first defined the Global Cambridge Investment Zone as a structure for retaining IP, talent, and ownership in Cambridge, and the second set out a model for producing startups more systematically within that structure. This paper asks how that zone can itself scale as a coordinated economic system. It argues that Cambridge must develop not only stronger internal coordination, but also durable commercial relationships with peer hubs elsewhere, so that the ecosystem operates as an active participant in global trade rather than a passive source of knowledge. The paper proposes that this can be achieved through automated coordination infrastructure, embedded financial and insurance systems, cross-hub dealmaking, and the development of AI capabilities shaped by the ecosystem's own operational needs. In this view, the GCIZ becomes both a mechanism for venture growth and a platform for producing exportable services in finance, risk, and AI-enabled coordination. Together, the three papers set out a progression from zone formation, to venture production, to ecosystem scale.

Contents

1	Strategic Overview	5
2	Framing: The Operating System	6
2.1	Core Concept	6
3	Context: Britain’s AI Gap and Cambridge’s Position	12
3.1	The Gap That Matters	12
3.2	The State of the Ecosystem	13
3.3	The Rise of Ecosystem Business Models	13
3.4	From Agentic Firms to Agentic Ecosystems	14
3.5	Agentic AI: Production Readiness and the Execution Gap	16
3.6	Geopolitical AI Sovereignty	16
3.7	InsurTech and Parametric Finance	17
4	Problems: From Production to Scaling	17
5	Change Justification: From Production Model to Operating System	19
5.1	Defining Scaling — A Gap in the Literature	19
5.2	A Triple Definition of Scaling	20
5.3	The Five-Stage Growth Pathway	20
5.4	What Scaling Enables	21
5.5	Stakeholders by Stage	21
6	Framework Principles	22
6.1	Closing the Scaling Loop	24
7	FinTech as the Mechanism for Startup Scalability	25
8	Global Trade and Supply Chain Risk Management	26
9	Intervention Network Analysis	27
10	Proposals by Layer	27
10.1	Layer 1: Ecosystem Sensing	27
10.1.1	1a. Live Venture Registry (Stages 1–2).	27
10.1.2	1b. Needs-Offers Matching Engine (Stages 2–3).	27
10.1.3	1c. Network Health Diagnostics (Stages 3–4).	28
10.1.4	1d. Sector Emergence Detection (Stage 4).	28
10.1.5	Scaling value.	28
10.2	Layer 2: Operational Automation	28
10.2.1	2a. Growth-Stage Governance Routing (Stages 1–3).	29
10.2.2	2b. Contract Generation and Procurement Compression (Stages 2–3).	29

10.2.3	2c. Financial Operations Automation (Stage 2).	29
10.2.4	Scaling value.	29
10.3	Layer 3: Intelligent Capital Coordination	30
10.3.1	3a. Proactive Deal Flow Matching (Stages 1–3).	30
10.3.2	3b. Dynamic Fund Structuring and Tax Arbitrage Optimisation (Stages 2–4).	30
10.3.3	3c. IP Valuation Engine (Stages 3–5).	30
10.3.4	3d. Value Retention Instruments (Stages 3–4).	31
10.3.5	3e. Parametric Innovation Insurance (Stages 2–5).	31
10.3.6	3f. Cross-Hub Co-Investment Vehicles (Stages 1, 5).	31
10.3.7	Scaling value.	31
10.4	Layer 4: Ecosystem-Level Dealmaking (Human-Led, AI-Supported)	32
10.4.1	4a. Hub-to-Hub Strategic Partnerships (Stage 5).	32
10.4.2	4b. Institutional and Policy Campaigns (Stage 5).	33
10.4.3	4c. Anchor Tenant Recruitment (Stages 1, 5).	33
10.4.4	4d. Geopolitical and Trade Intelligence (Stages 4–5).	33
10.4.5	Scaling value.	34
10.5	Layer 5: Financial Self-Sustainability	34
10.5.1	5a. Membership Tiers — Community and Full, with Mutual Stake (Stages 1–5).	34
10.5.2	5b. Compliance-as-a-Service (Stages 2–5).	35
10.5.3	5c. Transaction Fees on Capital Coordination (Stages 3–5).	35
10.5.4	5d. Parametric Insurance Premium Income (Stages 3–5).	35
10.5.5	5e. Data Products — Anonymised Ecosystem Intelligence (Stages 3–5).	36
10.5.6	5f. IP Valuation Tool Licensing (Stages 4–5).	36
10.5.7	5g. SaaS Licensing of Ecosystem OS (Stage 5).	37
10.5.8	Scaling value.	37
10.6	Layer 6: AI Ecosystem Generation	37
10.6.1	6a. Training Data Creation (Stage 3 onward).	38
10.6.2	6b. Demand Pull on Cambridge AI Research (Stages 3–4).	38
10.6.3	6c. Controlled Deployment Environments (Stage 4).	38
10.6.4	6d. Ecosystem Control Mechanism (Stage 4).	39
10.6.5	6e. The Innovation Prediction Model — Moonshot Output (Stages 4–5).	39
10.6.6	6f. Open-Source Release Strategy (Stage 5).	39
10.6.7	Scaling value.	40

11 Conclusions 40

11.1	Alternative Approaches and Cambridge’s Unique Conditions	40
11.2	Relationship to Prior Global Cambridge Models	41
11.3	Frameworks developed	42
11.4	Key Observations	43

11.5 Next Steps 44

11.5.1 Actions for GCIZ Stakeholders (Operators, Local Institutions, Investors) 44

11.5.2 Actions for Governmental Entities (Local, National, Regulatory) 46

11.5.3 Actions for Partner Hubs (International and Domestic) 47

11.6 Next Research Steps 48

1 Strategic Overview

Cambridge produces an exceptionally high concentration of world-class scientific research relative to the investment it attracts, even though it is already one of the most active venture capital ecosystems in the world on a per capita basis. The problem is not a lack of ideas, talent, or technical quality. It is that too little of this scientific output is retained, coordinated, and scaled in ways that allow Cambridge to capture more of the long-term value it creates.

This is the third paper in a series by the Global Cambridge Institute. The first, the *GCI Strategy Plan* (28), set out the central paradox of the Cambridge economy: exceptional research strength alongside structurally weak retention of IP, talent, and ownership. It proposed the Global Cambridge Investment Zone (GCIZ) as a practical response: a government-owned, GCI-operated framework combining tax advantages, curated membership, shared services, and public dashboards to help firms remain and grow in Cambridge. The second paper, *Systematising Startup Creation* (44), addressed how startups could be created and supported more systematically within that environment, proposing the Cambridge Startup Production Model — a coordinated system linking training, capital, production, currency, and culture to make venture creation more reliable and more deeply embedded in the city.

This paper asks how the zone itself can scale. Where the first paper defined the zone and the second explained how ventures could be produced within it, this paper treats the ecosystem as an organised economic entity in its own right. Cambridge must become more than a source of knowledge for others to extract. It must coordinate its own assets, form durable commercial relationships with other hubs, and operate at a level above the individual firm. That requires three things at once: growth that does not proportionally increase coordination complexity; stronger deal relationships with peer hubs, so that scale is achieved by becoming part of a wider network of interoperable innovation systems; and these two elements must amplify one another — autonomy enabling hub connectivity, and hub connectivity enabling autonomy.

The infrastructure through which this is achieved has three central components. Agentic AI must automate administrative, financial, legal, and matching functions that would otherwise remain fragmented and labour-intensive. Financial technology — understood not as one sector among others but as the underlying coordination layer — is equally central, since every serious act of organisation, exchange, risk allocation, and growth eventually has a financial expression. Given present geopolitical conditions, this extends to InsurTech: insurance and trade-risk systems are particularly relevant where cross-border operations, supply chain resilience, and new market formation are concerned.

Britain is unlikely to dominate the hardware and platform layers of AI in the manner of the United States or China. Its stronger position lies in service infrastructure: in the design of systems that coordinate finance, law, insurance, trade, and business operations. The GCIZ is therefore relevant not only as a local development tool but as a testbed for exportable forms of AI-enabled service infrastructure. Cambridge’s existing strengths in supply chain finance, risk

management, and advanced research can be operationalised in ways that generate both local retention and wider commercial value.

This paper also argues that the data generated by such an ecosystem is not merely a by-product, but a further enabling mechanism of autonomy – it is one of the key strategic assets produced by the system itself. This concept derives itself from principles of circular economy. A trade-connected, operationally active innovation zone generates the kinds of real coordination data from which more capable and better-grounded AI systems can be built. In that sense, the ecosystem is not only using AI to scale; it is also creating the conditions under which new AI capabilities can be developed and governed. This has implications for sovereignty as well as commerce. An ecosystem-level AI architecture that is community-owned, transparently governed, and not dependent on the strategic interests of either the United States or China can become attractive to other hubs seeking greater autonomy over their technological and economic infrastructure.

Taken together, the three papers now define a full growth pathway for the Cambridge innovation ecosystem: **Zone, Production, Scale**. The first establishes the zone. The second explains how ventures can be produced within it. This third paper explains how the system itself can grow, connect outward, and compound its advantages over time.




Figure 1: Paper Series Progression — Zone, Production, Scale

Readers unfamiliar with the first two papers should consult them for the fuller diagnosis of the Cambridge ecosystem and the detailed design of the earlier interventions. This paper recaps the essential logic, but its purpose is different: to set out how the GCIZ can move from being a local retention mechanism to becoming a scalable economic system, capable of coordinating growth within Cambridge and building productive relationships across a wider network of innovation hubs.

2 Framing: The Operating System

2.1 Core Concept

Systematising Startup Creation (44) established that Cambridge’s ecosystem — ranked first in the UK for science and technology intensity, with 848 active companies (6) generating nearly

£30 billion annually (56) — loses late-stage value to overseas acquirers while coordinating internally through emails, spreadsheets, and personal introductions. The Production Model proposed there addresses what to build; this paper addresses how to run it at scale.

The paper sets out a five-stage growth path for six infrastructure layers and 29 specific interventions. Each stage has measurable thresholds that must be met before advancing.

The growth path is as follows:

- **Stage 1 — Establish:** Import external tools, stand up initial infrastructure, begin data collection. *Threshold:* at least 3 external tool partnerships active, data flowing into a unified registry, initial membership tier operational.
- **Stage 2 — Naturalise:** Adapt imported tools to Cambridge’s specific context (UK tax regime, SEIS/EIS, university IP). *Threshold:* naturalised tools demonstrating at least 50% cost/time reduction in at least one coordination function.
- **Stage 3 — Integrate:** Achieve system coherence across layers so they feed each other end-to-end. *Threshold:* at least 3 function classes handled end-to-end, the data flywheel generating labelled training sets, at least one revenue stream financially sustainable.
- **Stage 4 — Generate:** Produce Cambridge-native AI models and technologies that no other hub can replicate. *Threshold:* at least one Cambridge-native AI model in production, ecosystem control mechanism operational, at least 2 partner hubs formally interested.
- **Stage 5 — Export:** License the ecosystem OS globally, generating revenue and inter-hub leverage that feeds back to Stage 1.




Figure 2: Five-Stage Growth Pathway

These stages each represent a step through the phases of scaling of a given ecosystem. This is used to guide and gauge how the following infrastructural layers are grown:

- **Layer 1 — Sensing:** Mechanisms that deliver ecosystem sensing capabilities. A live venture registry, needs-and-offers matching, network diagnostics, and tools to identify emerging sectors.
- **Layer 2 — Operations:** Tools to enable operational automation. Governance routing, contract generation, faster procurement, and automation of financial operations.

- **Layer 3 — Capital:** Means of intelligent capital coordination. Deal-flow matching, fund structuring, IP valuation, value-retention instruments, parametric insurance, and cross-hub co-investment.
- **Layer 4 — Dealmaking:** Connections to permit ecosystem-level dealmaking. Strategic action led by people but supported by AI, including hub partnerships, institutional campaigns, anchor recruitment, and geopolitical intelligence.
- **Layer 5 — Financial Sustainability:** Methods to ensure the system is self-sustainable. Revenue generated across all layers through SaaS licensing, compliance-as-a-service, IP valuation tools, data products, transaction fees, insurance premiums, and tiered membership.
- **Layer 6 — AI Generation:** Tools to initiate, mechanistically, the creation of novel AI ecosystems. Training data creation, demand-pull for AI research, and ecosystem control mechanisms. Its most ambitious element is an innovation prediction and development model trained on longitudinal GCIZ data.




Figure 3: Six Infrastructure Layers with Data Flywheel

The layers are connected through a data flywheel: each layer’s operations generate data that trains AI that improves every layer, working as both the mechanism of scaling and the system’s primary export. They are broken down into the collection of interventions in Table 1.

Intervention	Stages	Core function
--------------	--------	---------------

1a. Live Venture Registry	1–2	Ingest open access company and innovation data, including company listings, patents, funding, hiring signals
1b. Needs-Offers Matching Engine	2–3	Live index of what every venture needs and offers; replaces manual brokering
1c. Network Health Diagnostics	3–4	Centrality measures, structural hole detection, digital twin of the ecosystem
1d. Sector Emergence Detection	4	Patent co-filing, hiring clusters, co-authorship signals to break deeptech/medtech lock-in
2a. Governance Routing	1–3	Automating venture governance for things like regulations, IP, data, and finances through inputted company profiles
2b. Contract Generation & Procurement	2–3	AI-drafted NDAs, IP assignments, investment terms to optimise partnership creation
2c. Financial Operations Automation	2	Invoicing, grant reporting, cap tables, SEIS/EIS claims, R&D tax credits
3a. Proactive Deal Flow Matching	1–3	Continuous matching of investor mandates against live registry
3b. Dynamic Fund Structuring	2–4	Model optimal vehicles given SEIS/EIS, investor profiles, deal flow
3c. IP Valuation Engine	3–5	Train on longitudinal Cambridge data to produce defensible deep-tech IP valuations, simplifying dealmaking
3d. Value Retention Instruments	3–4	Convertible preferences, trailing royalties, ecosystem equity pools to retain stakeholdership through exits
3e. Parametric Innovation Insurance	2–5	AI-driven risk assessment with trigger-based payouts (regulatory change, key person loss, IP challenge, trade risk)
3f. Cross-Hub Co-Investment Vehicles	1, 5	Multi-jurisdictional vehicles with automated cross-border compliance

4a. Hub-to-Hub Strategic Partnerships	5	Bilateral agreements with other innovation hubs, such as London, Dubai, Boston, Singapore, Tel Aviv
4b. Institutional & Policy Campaigns	5	Lobby for expanded SEIS/EIS, establish more educational pathways to entrepreneurship, engage innovation funding bodies
4c. Anchor Tenant Recruitment	1, 5	Attract corporates, sovereign wealth funds, and international research organisations
4d. Geopolitical & Trade Intelligence	4–5	AI-mediated analysis of trade restrictions, sanctions, sovereignty movements
5a. Cooperative Governance	1–5	Cooperative entity with tiered membership and mutual stake; community and full tiers with differentiated access
5b. Compliance-as-a-Service	2–5	Productised governance routing and contract generation as subscription service
5c. Transaction Fees on Capital	3–5	To finance growth, a cost capture mechanism on capital deployed within matching and structuring infrastructure
5d. Parametric Insurance Premiums	3–5	Premium income from innovation insurance products; counter-cyclical revenue buffer
5e. Data Products	3–5	Anonymised ecosystem intelligence: venture trajectories, deal outcomes, sector signals
5f. IP Valuation Tool Licensing	4–5	Licensing the IP valuation engine to investors, rating agencies, and pension funds globally
5g. Ecosystem Licensing	5	Licensing the full integrated platform to other innovation hubs as configurable SaaS
6a. Training Data Creation	3+	Aggregate labelled datasets from all layers (venture trajectories, deal outcomes, governance mappings)

6b. Demand Pull on AI Research	3–4	Unsolved framework problems become applied research opportunities for Cambridge AI groups
6c. Controlled Deployment Environments	4	Bounded autonomy architectures with human escalation and audit trails — a production testbed for agentic AI
6d. Ecosystem Control Mechanism	4	Bespoke AI for managing the integrated ecosystem in real time — the Cambridge-native product
6e. Innovation Prediction Model	4–5	The “moonshot” — trained on 3–5 years of longitudinal ecosystem data to predict outcome distributions

Table 1: Interventions

We define an **agentic innovation ecosystem** as a geographically anchored but globally connected network of autonomous entities (ventures, investors, universities, service providers) whose routine coordination is performed by purpose-built AI agents operating within cooperative governance structures, sustained by embedded financial infrastructure, and producing foundational AI as both a by-product of operation and a primary export.

This extends the concept of the agentic organisation (54) beyond the single firm to the harder inter-firm problem of ecosystems, where no singular authority exists and participants can enter and exit at will. We leverage cooperative governance following empirically validated commons principles (42) rather than corporate hierarchy. The architecture is practically implemented following the notion of the Agentic Mesh (7), adding domain specificity and financial self-sustainability. A related case appears in the Lutech Campus case (5), but the model proposed here differs in scale (city-level versus a single hub), governance (cooperative rather than corporate), and the dual role of AI as both instrument and product.

The paper’s five novel contributions are:

1. A triple definition of ecosystem scaling — compression, amplification, resilience — filling a gap across the entrepreneurial ecosystem, regional innovation system, and platform literatures
2. A five-stage growth pathway as the mechanism of scaling
3. The agentic innovation ecosystem as a new organisational category
4. Fintech and InsurTech as organising substrate rather than revenue layer
5. A data flywheel that is simultaneously the mechanism of progression through growth stages and the system’s primary output, culminating in an innovation prediction model

trained on data no other hub possesses




Figure 4: The Five Key Contributions of the Paper

3 Context: Britain's AI Gap and Cambridge's Position

3.1 The Gap That Matters

Britain does not lack AI research capability. Cambridge alone hosts a founder-to-researcher ratio of 1:12 (versus 1:31 in London) (17). The UK is third globally for innovation output and fifth overall for innovation (58). What Britain lacks is agentic AI infrastructure for its service economy. The US and China dominate AI at the platform and hardware layers — foundation models, cloud infrastructure, chip fabrication. Britain cannot compete there and should not try. Britain's competitive advantage is its service economy — financial services, insurance, legal, and professional services — constituting over 80% of GDP and making it one of the largest services exporters in the G7 (32).

The AI gap is in the agentic systems that automate, coordinate, and scale that service economy: InsurTech (parametric risk products, AI-driven underwriting, automated claims), FinTech (deal flow automation, compliance routing, fund structuring, cross-border payment rails) and LegalTech (contract generation, IP governance, regulatory navigation). These are domains where Cambridge has world-class research — CCAF's fintech dataset covers 2,000+ entities across 14 market segments; the Judge Business School runs global fintech executive programmes; Cambridge's supply chain finance research at the Institute for Manufacturing is internationally recognised — yet conspicuously underproduces companies. The proposed GCIZ is a vehicle through which Cambridge can close this gap: not by pivoting away from deeptech and life

sciences, but by recognising that FinTech, InsurTech, and LegalTech are the scaling infrastructure that deeptech and life sciences need to reach global markets.

CCAF’s Cambridge Fintech Ecosystem Atlas tracks 2,000+ entities across 14 market segments (11); Cambridge Judge Business School offers executive fintech and regulatory innovation programmes with an explicitly global policy and decision-maker focus (14); and Cambridge’s Institute for Manufacturing has developed internationally used supply-chain research applied in more than 35 leading multinationals (51). Despite the impactful and respected research at Cambridge, only a small fraction of its research and IP activity converts into new companies: Cambridge Enterprise reported 25 new companies formed in 2023–24 alongside 469 patent applications (12). The GCIZ is a vehicle through which Cambridge can close this gap: not by pivoting away from deeptech and life sciences, but by recognising that FinTech, InsurTech, and LegalTech are part of the scaling infrastructure those sectors need to reach global markets.

3.2 The State of the Ecosystem

Cambridge’s ecosystem has grown to 848 active companies, with innovation activity up almost 80% over the past decade. Investment intensity has accelerated sharply — 359% growth in equity per company and a 699% rise in annual funding since 2015 (6). The ecosystem now generates \$17.7 in startup value per VC dollar, outpacing both the Bay Area and Zurich. S-curve modelling suggests it could reach \$500 billion in value by 2030 if current bottlenecks are addressed (17).

Yet approximately 75% of knowledge-intensive turnover in Cambridge city is in foreign-owned companies (9). ARM — a Cambridge-headquartered chip designer — chose the NASDAQ for its 2023 IPO, listing at an initial valuation of approximately \$54.5 billion, the largest US listing since 2021 (36). The outputs of the Cambridge tech ecosystem are often ultimately listed or controlled overseas; in this sense Cambridge functions as an R&D engine for foreign owners and markets (34).

Cambridge’s fintech dimension is conspicuously underrepresented relative to its deeptech and life sciences strengths. The CCAF at Judge Business School runs one of the world’s largest fintech datasets (over 2,800 entities across 14 market segments (10)), yet Cambridge’s own ecosystem reports and Dealroom’s original analysis categorise the city primarily as a deep tech and life sciences hub, with no dedicated fintech sector breakdown (20). As this paper argues, financial infrastructure is the organising substrate of ecosystem scaling itself.

3.3 The Rise of Ecosystem Business Models

The business model landscape is shifting from linear product and service delivery to ecosystem orchestration. A 2025 survey of 700+ CEOs found 77% are rebuilding business models around AI, predicting it as the single biggest disruptor by 2028, yet only 35% of digital transformation initiatives realise intended outcomes. Platform business models have evolved from first-generation

multi-sided marketplaces into expansive cross-sector ecosystems that aggregate capabilities from partners, customers, regulators, and competitors. Value now arises from data synergies and co-innovation, not just scale (24).

Ecosystem businesses take fundamentally different organisational forms from traditional companies. A distinction exists between platform-based ecosystems — where digital infrastructure mediates all interactions (Apple’s App Store, Stripe’s payment APIs) — and non-platform ecosystems, where cross-organisational collaboration generates value without a single technological substrate (49). Multiple distinct ecosystem business models exist, ranging from symbiotic (complementary offerings through a shared platform) to coopetitive (competitors cooperating to create higher combined customer value) (48). The critical organisational questions for any innovation ecosystem are: who orchestrates, and how are things coordinated? Options include corporate platforms, cooperative mutuals, foundation models (as with the Linux Foundation or Apache Foundation), and Decentralised Autonomous Organisation (DAO)-like structures where governance, rather than hierarchy, is encoded in rules.

For innovation ecosystems, the cooperative mutual form has strong appeal. It avoids the profit-extraction logic that misaligns traditional company structures with ecosystem-level goals, enables capital sharing while maintaining participant independence, and can be structured as tax dormant. The UK’s Co-operative and Community Benefit Society Act 2014 provides a ready legal vehicle. For an illustrative example, we can use, “The Venture Society”, who demonstrate this model in practice: they function as a skill and experience-sharing collective for entrepreneurs and fund managers, with FCA registration enabling postnominal issuance as a credibility and culture layer, and a two-tier structure of community members at £10 (300 members) and full members pledging £10,000 (78 members), each contributing to a community fund (25). Such numbers are obtained upon manifestation and operation of the cooperative mutual entity so the precise values for a Cambridge based entity are not written here, however, they may be used as a benchmark.

Ecosystem payment architecture must balance accessibility with commitment alignment. Transaction-based models (Stripe’s percentage per transaction), commission-based models (Apple’s App Store 30%), carry-based models (AngelList syndicates), and tiered membership are each suited to different ecosystem types. Tiered membership is particularly relevant for innovation ecosystems because it creates mutual stake — participants who have invested in the ecosystem’s success have incentive to contribute, not merely extract. Technical implementation requires identity and membership management, payment rails (direct debit for subscriptions, escrow for mutual funds, parametric triggers for conditional disbursements), and transparent ledger systems for contribution tracking.

3.4 From Agentic Firms to Agentic Ecosystems

The concept of layered autonomy in distributed networks has been developed across several disciplines. Notable are the manifestations in logistics and supply chain: (59; 60) propose a

seven-level autonomy reference for autonomous supply chains, and how to implement autonomy. (18) argue supply networks are complex adaptive systems where “imposing too much control detracts from innovation and flexibility; conversely, allowing too much emergence can undermine managerial predictability”. (33) formalise autonomous cooperation in logistics. Each demonstrates that distributed networks of independent entities can self-govern across organisational boundaries when given appropriate structure.

At firm level, (54) proposes “the agentic organisation” — a paradigm uniting humans and AI agents to work side by side at near-zero marginal cost. Organisation charts pivot from hierarchical delegation toward agentic networks based on exchanging tasks and outcomes. Critically, McKinsey observes that “agentic networks are not necessarily limited to the boundaries of a single organisation, and different outcomes may be sourced from different parties, opening up new B2B opportunities”. If the cost of each additional coordination task approaches zero once infrastructure exists, the system can grow without proportional cost increase — the quantitative foundation of scaling as compression.

(7) extends this further with the concept of an Agentic Mesh — an interconnected ecosystem where federated autonomous agents and people initiate and complete work together, designed to support agent collaboration, foster trust, and enable agents to find each other, safely collaborate, interact, and transact. The Lutech Campus case study provides the closest existing empirical example: an AI-based open innovation hub facilitating knowledge sharing, decision-making, and collaborative development across companies, startups, and academic institutions (5).

The **agentic ecosystem** is a novel concept beyond the agentic firm. An agentic organisation has a single entity setting governance, defining agent roles, controlling data, and bearing accountability. An agentic ecosystem has no single governing entity — participants are autonomous, with independent decision-making capability. Agents serve different masters. Governance must be distributed, cooperative, and designed for entities that can exit at will. This is fundamentally harder. The framework addresses it by combining layered autonomy from distributed networks literature, cooperative commons governance from Ostrom, fintech and InsurTech as financial substrate, and purpose-built AI coordinating this together.




Figure 5: From Agentic Organisation to Agentic Ecosystem

3.5 Agentic AI: Production Readiness and the Execution Gap

The global AI agents market reached approximately \$7.6–7.8B in 2025, projected to exceed \$50B in 2026 and \$52 billion by 2030 (29). (26) forecasts 40% of enterprise applications will embed task-specific AI agents by end of 2026, up from under 5% in 2025. (47) surveyed US business leaders, and show 79% of organisations report some AI agent adoption, with 88% planning expansion of AI use.

Conversely, (21) finds that over 40% of agentic AI projects are expected to fail by 2027 due to organisation problems. (46) likewise shows that only a minority of organisations have progressed to more advanced stages of AI governance maturity, with many still building basic structures, training, and controls. (21) further identifies three fundamental infrastructure obstacles to scaling agentic AI: legacy system integration (enterprise systems not designed for agentic interactions), data architectures that limit searchability and reusability, and immature governance frameworks for autonomous systems.

Cambridge has a structural advantage on all three. Cambridge is built on flat, decentralised operation, thus there is no legacy ecosystem infrastructure to integrate. Also, its AI and computer scientific competency allows it to use cutting edge data management techniques on itself for agentic use. In a sense, it already has governance built on Ostrom’s principles, positioning it perfectly for scaling infrastructure. The OECD’s finding that AI lowers minimum efficient scale (40) validates Cambridge’s SME structure as a competitive advantage when paired with the right infrastructure. We propose a five-stage pathway (defined earlier) to enable this and avoid the “perpetual pilot trap”. This is ensured by each stage being phased with financial sustainability requirements and measurable outcomes. Nothing proceeds without demonstrated value.

3.6 Geopolitical AI Sovereignty

The AI stack is increasingly controlled by a small number of US and Chinese players, creating risks related to data access, model restrictiveness, and vertical integration (40; 39). For innovation ecosystems, dependency on foreign AI providers creates structural vulnerability: pricing changes, access restrictions, terms of service modifications, or compliance with home-government geopolitical imperatives, none of which the ecosystem controls.

Cambridge is exceptionally well positioned as a UK AI hub, particularly in research-led and deeptech AI: the ai@cam regional statement notes that the Cambridge region contains 120 AI-powered companies employing 13,000 people with a combined turnover of £6 billion, alongside 36 research parks, two universities, and a strong startup and investor base (1). This is reinforced by the University of Cambridge’s continuing pipeline of AI spinouts through Cambridge Enterprise (55) and by the UK government’s January 2026 announcement of £36 million to expand Cambridge’s AI Research Resource capacity sixfold, building on DAWN as one of the UK’s most powerful AI supercomputers (23).

The UK government’s AI Sector Study highlights that the UK has a strong AI ecosystem, particularly at the early stages of company formation, with clear scope to further leverage this vibrant research, innovation and start-up activity through higher-value, later-stage capital investment (22). Cambridge sits at the centre of this opportunity: as the UK’s most research-intensive AI cluster, it is uniquely placed to convert upstream strength — world-class talent, research, and a deep early-stage pipeline — into the next generation of globally significant AI companies built and scaled in the UK, with the AIRR expansion and the wider Oxford–Cambridge corridor providing precisely the compute, capital and infrastructure substrate that scaling requires.

This makes an OS for the ecosystem attractive not only as a productivity tool but as a geopolitically non-aligned alternative to major powers. Ecosystem data stays within cooperative governance; models are community-owned; deployment requires no foreign infrastructure. Any hub concerned about technology sovereignty — in the EU, Singapore, the Gulf states, India, etc. — has interest in ecosystem management infrastructure free from dependency on a single foreign AI provider.

3.7 InsurTech and Parametric Finance

Parametric insurance — payouts triggered by predefined measurable events (weather thresholds, seismic data, vessel tracking, port congestion) rather than loss assessment — is maturing rapidly. The global protection gap is projected at \$1.86 trillion (45). The wider AI-native insurance space is also seeing significant investment: Corgi Insurance secured \$108 million and full regulatory authority in early 2026 to operate as an AI-native, full-stack insurance carrier built specifically for startups (16). A new generation of insurtechs more specifically focused on parametric products and decentralised risk pools is also emerging (16). LLMs are becoming distribution’s next competitive battleground in insurance, with agentic AI reshaping both embedded and traditional channels (16).

Supply chain disruption — from COVID-19, the Suez Canal blockage, semiconductor shortages, sanctions regimes, and geopolitical realignment — has exposed the fragility of trade networks. Supply chain risk is no longer merely operational; it is strategic, financial, and insurable. Marine cargo parametric products (triggered by vessel tracking, port congestion metrics, weather data) and commodity-backed property instruments (triggered by price indices, trade flow disruption) represent direct applications for Cambridge’s trade finance research base.

4 Problems: From Production to Scaling

Systematising Startup Creation (44) identified nine interconnected problems at the venture-creation level. The interventions proposed there — Training, Capital, Production, Currency, and Culture layers — address these at that level. A further set of problems emerges when the question shifts from "how do we create more ventures?" to "how do we scale the ecosystem that

creates them?" These are:

- 1. Distribution deficit.** Cambridge builds world-class technology but lacks routes to market. Banks have centuries of KYC'd client bases; when tokenisation platforms or novel financial infrastructure products complete development, they face no equivalent customer network. Across blockchain infrastructure, compliance tools, and ecosystem management itself, the binding constraint is shipping, not building.
- 2. Coordination deficit.** The previous paper documented Cambridge's fragmentation — bottom-up governance, siloed departments, weak cross-sector ties — and proposed connective tissue through the Production Model. Coordination at the scale of 848 companies, 36 research parks, and 5,000+ innovation-driven firms (35) cannot be addressed through mechanisms that work for a community of a few hundred active participants. Knowledge lives in individuals' heads and meeting notes. London LPs have been unable to penetrate Cambridge investment networks — not because deals don't exist but because finding them costs more than investors will tolerate. The sector lock-in (deeptech and medtech dominating, fintech underrepresented despite world-class research) reflects path dependency, not capability absence.
- 3. Late-stage capital flight and extractive international linkages.** Foreign capital accounts for 60% of late-stage VC investments in the UK, versus 15% in the US. US corporates and funds acquired nearly \$24 billion of European spinout value since 2019 (3). UK-based investors accounted for 93% of Cambridge deals in 2015 but just 60% in 2024, with US participation more than doubling from 8% to nearly 19% of all rounds (6). Haskell and Westlake estimate UK underinvestment in its innovation economy could cost up to £143 billion annually in lost potential (31). Cambridge-born companies — Darktrace (acquired by Thoma Bravo), Featurespace (acquired by Visa for £700M), IQGeo (acquired by KKR for £333M), and ARM (SoftBank, \$31B) — represent billions in value transferring to overseas ownership. CIC launched its £100M Opportunity Fund explicitly to address the point where "the UK often loses its most exciting businesses" (13). Cambridge's international relationships are presently extractive: the ecosystem functions as an R&D farm for overseas acquirers, with international linkages structured as buyer-seller acquisitions rather than peer-to-peer co-investment, co-development, or strategic partnership.
- 4. Governance deficit.** The governance challenge generalises beyond regulatory compliance to IP governance (who owns what when a spinout is formed), data governance (how ecosystem intelligence is shared and protected), financial governance (fund structures, tax incentive navigation, cross-jurisdictional investment), and institutional governance (university rules, FCA requirements, cooperative law). No scalable, integrated governance navigation layer currently exists. Cambridge Enterprise's IP terms — while comparatively favourable among UK universities — still push businesses toward early acquisition rather than local scaling.
- 5. Procurement cycle drag.** Corporate innovation procurement takes 6–18 months. Open innovation is blocked not by technology readiness but by purchasing process, which directly suppresses scaling velocity.

6. Strategic misorientation. AI has shifted the scarce resource from operational execution to strategic partnership design. Treating scaling as an operational challenge — more events, more introductions, more administration — rather than a strategic one — which partners, which structures, which instruments — is the wrong approach. Among organisations with extensive agentic AI, 45% expect to flatten management hierarchies (39). The question is no longer what skills are needed, but what the structure should look like when humans and agents work side by side.

7. FinTech underrepresentation. CIC's investment portfolio targets "AI, IoT, quantum, autonomous systems, therapeutics, medtech, diagnostics, genomics" — fintech is absent (13). Globally, fintech attracts approximately 10.9% of new VC fund allocations, yet in Cambridge the sector is not even measured as a category in the city's primary ecosystem reports. Featurespace — Cambridge's most prominent financial technology success — was acquired by Visa rather than scaling locally. The gap matters directly for scaling: fintech provides the infrastructure tools (deal matching, compliance automation, fund structuring, insurance products, payment rails, identity systems) that the ecosystem itself needs to scale. Cambridge researches these tools globally but does not build them locally.

5 Change Justification: From Production Model to Operating System

5.1 Defining Scaling — A Gap in the Literature

The innovation ecosystem literature uses "scaling" pervasively but defines it loosely. The entrepreneurial ecosystem tradition (37; 52; 53) treats scaling as the growth of individual ventures within an ecosystem — the ecosystem is context, not the object of scaling. The regional innovation system tradition (19; 4) identifies configurations that produce high-tech outcomes but does not operationalise the scaling of ecosystem infrastructure itself. The platform tradition (27; 2) treats ecosystems as central to value creation; (2) in particular defines an ecosystem as "the alignment structure of the multilateral set of partners that need to interact in order for a focal value proposition to materialize," but the mechanism by which that structure scales without proportional coordination cost increase is unspecified. Granstrand and Holgersson (30) synthesise innovation ecosystem definitions as "the evolving set of actors, activities, and artefacts, and the institutions and relations" — scaling is implicit in "evolving" but never operationalised. The OECD/IDIA framework (41) emphasises analysing ecosystem maturity levels — a state description, not a mechanism. The AI-enabled scaling literature (57; 50) shows AI enabling firm-level ecosystem scale-up but does not address multi-stakeholder, geographically distributed innovation ecosystem scaling.

The gap is threefold: no existing definition addresses what changes when an ecosystem scales (not just that "it gets bigger"), what infrastructure is required to enable scaling without proportional

cost increase, and what mechanism converts infrastructure into scaling outcomes.

5.2 A Triple Definition of Scaling

Compression. This reduces the unit cost of each ecosystem coordination function. McKinsey’s near-zero marginal cost framing for agentic systems provides the quantitative foundation (54): once infrastructure exists, each additional coordination task approaches zero marginal cost. This is measurable as time and cost per coordination event.

Amplification. This enables qualitatively new operations impossible at craft scale. One-to-one business development becomes ecosystem-to-ecosystem connections. This is measurable as the number and type of novel function classes the ecosystem can perform.

Resilience. This eliminates single points of failure. Knowledge must be decentralised, pathways multiplied, and no individual gatekeeper critical. This is measurable as recovery time and the inverse of capability loss when key nodes are removed.



Figure 6: Triple Definition of Scaling. Scaling = Compression + Amplification + Resilience

5.3 The Five-Stage Growth Pathway

The five-stage growth pathway is the mechanism by which this scaling is achieved, with each stage producing measurable advances on all three dimensions:

1. **Establish Links** — initial compression via imported tools
2. **Naturalise** — deepened compression and early amplification via local adaptation
3. **Integrate** — amplification at scale via system coherence
4. **Generate** — all three dimensions via Cambridge-native technologies
5. **Export** — revenue and inter-hub leverage, feeding back to Stage 1

Stage transitions rely on the establishment of ecosystem at that tier. Stage 1→2 requires at least three external tool partnerships active, data flowing into a unified registry, and an initial membership tier operational, forming something beyond a simply dyad. Stage 2→3 requires naturalised tools demonstrating effective automation in time and cost for at least one coordination function, manifest between at least three distinct tool partnership clustering, with

Cambridge-specific regulatory mapping complete for SEIS/EIS and university IP. Stage 3→4 requires the integrated system handling at least three function classes end-to-end, the data flywheel generating labelled training sets, and financial sustainability demonstrated for at least one revenue stream across these cluster entities. Stage 4→5 requires at least one Cambridge-native AI model deployed in production, the ecosystem control mechanism operational, and at least two partner hubs expressing formal interest in adoption. For each, the task of progression is a repeatably functional layer that operates without active enforcement over more than a dyadic relationship, which serves as evidence of repeatability.

The data flywheel can stall through low participant contribution, poor data quality, or underperforming models. Mitigations to this must be built into its design: participant contribution must be incentivised through tiered membership where data-contributing members receive enhanced services (following the proportional equivalence concept (42)); data quality is maintained through automated validation at ingestion plus human curation at the sensing layer; model underperformance is addressed through the controlled deployment environment (Layer 6c), which provides structured feedback loops before any model enters production. The financial sustainability requirement at each stage gate ensures the flywheel is not dependent on continued external funding.

5.4 What Scaling Enables

With sufficient scalability, an ecosystem should become an exportable operating system that other hubs can adopt. Doing this for Cambridge would convert it from a single cluster with strong bridging of other hubs into the orchestrator of a global network of innovation ecosystems. Distributed and decentralised networks literature provides the design requirements: a layered model where each layer handles a distinct function class (33); defined autonomy levels specifying automation versus human intervention at each layer; structural entities bridging external and internal connections (59; 18); and interoperability across entities of varying sophistication. Routine coordination is automated at lower layers; strategic decisions remain human-in-the-loop at upper layers; entities from solo founders to multinational corporates interact through the same infrastructure regardless of individual scale.

5.5 Stakeholders by Stage

- **Stage 1 (Establish Links):** Cambridge Enterprise, Innovate Cambridge, CIC (13), Parkwalk, Cambridge Angels, CCAF, 2–3 target partner hubs (DIFC Dubai, EDB Singapore, BIRD Israel model), RegTech providers (ComplyAdvantage, Ascent).
- **Stage 2 (Naturalise):** Judge Business School, Cambridge law firms (HCR Hewitsons, Mills & Reeve), SEIS/EIS specialists, Accelerate Cambridge, university legal and IP offices, Carta/AngelList for financial operations adaptation, IfM supply chain finance group.
- **Stage 3 (Integrate):** CCAF (financial data integration), Cambridge AI research groups




Figure 7: Nine governing principles

(CSIT, Machine Learning Group), ARIA, Innovate UK, British Patient Capital.

- **Stage 4 (Generate):** Cambridge AI safety researchers, sovereign wealth fund innovation arms (Temasek, Mubadala, PIF), international investors, Centre for Science and Policy (CSaP), Pro-Vice-Chancellor for Innovation.
- **Stage 5 (Export):** Partner hub governments and innovation agencies, SaaS distribution channels, open-source developer communities, bilateral R&D foundation models (BIRD-equivalent structures), GCI international network, global trade finance networks.

6 Framework Principles

We follow a collection of core principles to define the framework with which we pursue scaling. We can see these principles summarised in Figure 7:

Principle 1: Automate Volume, Liberate Judgment. Each layer automates high-volume, rules-tractable functions, explicitly freeing human capacity for irreducibly human activities. Compliance routing is automatable; negotiating a cross-hub regulatory sandbox is not. Deal flow matching is automatable; deciding strategic partners is not. Lower layers operate at higher automation (routine, rules-based); upper layers remain human-led and AI-assisted (strategic, relational, political).

Principle 2: The Strategic-Operational Flip. AI moves the frontier from execution to ontological structuring. If defining what categories of venture exist, what relationships connect them, and what success metrics apply can be operationalised, ecosystems can be built as systematically as businesses. The framework operationalises ontology at lower layers so humans do strategy at upper layers.

Principle 3: Distribution Before Technology. For every component it must be clearly known who uses it and how they find it. Users experience services, not infrastructure. The ecosystem OS must be invisible to users at the point of use.

Principle 4: Governance as Competitive Advantage. Governance should be understood not as burden but as edge. Whoever navigates regulatory, IP, data, financial, and institutional governance more efficiently suffers fewer costs at the same targets. Automation optimises governance ROI per growth stage.

Principle 5: Fintech and Insurance as Organising Substrate. Fintech infrastructure should not be perceived as one sector among many, rather, as the medium through which every ecosystem relationship is formalised, executed, and sustained, such that social infrastructure is transformed into digital infrastructure, achieving autonomy. This includes introduction of membership tiers to formalise trust and engagement; deal flow matching and co-investment vehicles to formalise partnerships; transaction fees and insurance premiums to generate self-sustaining revenue; and smart contracts to provide identity and audit backbone on which all other coordination depends. Counter-intuitively, in the context of the other principles, this should optimise for the creation of fintech innovation.

Principle 6: Multi-Hub Connectivity as a Design Primitive. The framework is designed from inception for connectivity with peer hubs, with mind to drive network management and dealmaking to occur at the ecosystem. Separation on function is antithetical to the ecosystem scalability concept, however effective autonomisation of internal functions should permit strategisation for this. London for fintech innovation, Dubai for sovereign wealth and RAW tokenisation, Cambridge Massachusetts for knowledge exchange, Tel Aviv for R&D, Silicon Valley for capital, and emerging hubs such as Astana, Bangkok, and Lagos as viable testbeds for scalability concepts. Bidirectional interhub connectivity is the target outcome of scaling.

Principle 7: Portability as Dealmaking Mechanism. The system must work in other ecosystems — a design constraint and a business model simultaneously. When Cambridge seeds its OS in a partner hub, both operate on compatible architecture, enabling data exchange, co-investment, and coordinated governance. This is to follow the principle of nested enterprises, with local autonomy within shared protocols, scaling from hub to inter-hub to global network.

Principle 8: The Data Flywheel as Elevation Mechanism. The data flywheel is not merely a technical feature — it is the mechanism by which Cambridge elevates through growth stages while simultaneously producing its primary export product. At Stages 1–2, imported and naturalised tools begin generating data (venture trajectories, compliance pathways, deal outcomes). At Stage 3, integrated tools generate richer cross-layer correlations. At Stage 4, accumulated data enables training of models no other hub possesses the data to produce. At Stage 5, the models are the export product while also making the framework more effective locally, generating more data, and improving the models further. The flywheel is therefore both facilitator (driving Cambridge through growth stages) and goal (the AI models it produces are the system’s ultimate output).

Principle 9: Trust is Architectural The innovation ecosystem should be a common-pool resource: shared infrastructure (data, networks, reputation, institutional relationships) that participants can contribute to or extract from. To achieve this, we target the following properties:

- Clearly defined boundaries → tiered membership, defined ecosystem perimeter, clear data-sharing terms
- Proportional equivalence between benefits and costs → mutual stake tiers, contribution-

proportional access, equity-based value sharing

- Collective-choice arrangements → cooperative governance under the Co-operative and Community Benefit Society Act 2014, member voting, distributed decision-making
- Monitoring → AI-enabled transparent contribution tracking, automated ledgers, real-time ecosystem health dashboards
- Graduated sanctions → reduced access before expulsion, reputation impact before financial penalty
- Conflict resolution → built-in arbitration protocols, community mediation, transparent dispute records
- Minimal recognition of rights to organise → FCA registration, cooperative legal status, formal recognition by university and government
- Nested enterprises → local ecosystem governance nested within multi-hub agreements, venture autonomy within community norms, hub governance within inter-hub protocols

6.1 Closing the Scaling Loop

The framework explicitly closes the loop by connecting each output to the next input:

- Layer 1 Sensing: reveals needs and opportunities, feeding Layer 2's Operational Automation (what needs compliance? what contracts?) and Layer 3's Capital Coordination (what deals? what valuations?).
- Layer 2 Operational Automation: reduces friction, enabling faster Capital Coordination and Dealmaking in Layers 3 and 4.
- Layer 3 Capital Coordination: provides returns, justifying Layer 4's Dealmaking by funding Layer 5's Sustainability and creating deal flow.
- Layer 4 Dealmaking: opens new hubs, expanding Layer 1's Sensing scope.
- Layer 5 Sustainability: funds all layers and validates the model for export.
- Layer 6 AI Generation: improves every layer, produces exportable tools, attracts talent, and creates new dealmaking leverage.

Each cycle through the loop increases data available to AI, the accuracy of sensing, the efficiency of operations, the quality of capital matching, the leverage of dealmakers, and the revenue of the platform — compounding geometrically because each layer's improvement enhances every other layer's performance. This is represented by Figure 8.




Figure 8: Multilayer Scaling Interactions

7 FinTech as the Mechanism for Startup Scalability

Cambridge’s ecosystem is officially characterised through its two pillars — deeptech and life sciences — with fintech absent from CIC’s investment portfolio (13) and not measured as a category in primary ecosystem reports (6; 35). Featurespace — Cambridge’s most prominent financial technology success — was acquired by Visa for £700M rather than scaling locally. Globally, fintech attracts approximately 10.9% of new VC allocations and remains one of the top two funded sectors in most major ecosystems.

Fintech is not a sector alongside deeptech — it is the infrastructure through which all sectors scale:

- **Social infrastructure.** Membership tiers, mutual stakes, and cooperative governance are financial instruments formalising trust between ecosystem participants.
- **Business infrastructure.** Deal flow matching, co-investment vehicles, IP valuation engines, and contract automation are financial products that formalise partnerships and reduce transaction costs from months to days.
- **Monetary infrastructure.** Transaction fees, insurance premiums, SaaS licensing, and carry structures generate the revenue that makes the GCIZ financially self-sustaining, eliminating dependence on grants.
- **Technical infrastructure.** Payment rails, KYC/AML systems, smart contracts, parametric triggers, and transparent ledgers provide the identity, authentication, and audit backbone on which every other coordination function depends — including, at the most basic level, the venture registry’s identity layer.

Cambridge should encourage FinTech company formation explicitly, leveraging CCAF, Judge Business School, and its supply chain finance research base to produce the tools the GCIZ itself needs — and the tools that constitute exportable products for global trade infrastructure.




Figure 9: FinTech as Organising Substrate

8 Global Trade and Supply Chain Risk Management

The GCIZ is designed from inception for investable global trade — connecting Cambridge to partner hubs through compatible infrastructure enabling co-investment, shared governance, and coordinated market access. The global trade context makes this urgent. Supply chain disruption from COVID-19; blockages of the Suez Canal, Strait of Hormuz, and Bab al-Mandeb Strait; semiconductor shortages; sanctions regimes; and geopolitical realignment have exposed the fragility of international trade networks. Supply chain risk is no longer an operational concern; it is strategic and financial. Cambridge is uniquely positioned to address this through the GCIZ:

Supply chain finance research. The Institute for Manufacturing has produced internationally recognised research on supply chain risk, network science applied to trade finance, and the quantitative modelling of systemic risk in supplier networks. CCAF’s alternative finance datasets cover global trade finance instruments. This research base — currently academic — can be operationalised through the GCIZ into commercial products.

Parametric insurance for trade risk. AI-driven risk assessment with trigger-based payouts for measurable events: vessel tracking, port congestion metrics, commodity price indices, trade flow disruption. Marine cargo parametric products and commodity-backed property instruments are direct applications for Cambridge’s research. Lloyd’s of London began as merchants sharing risk in a coffee house, formalised into an institution; the GCIZ applies the same principle with AI-native infrastructure (16).

Cross-hub co-investment vehicles. Multi-jurisdictional investment structures with automated compliance routing — enabling bilateral capital flows between Cambridge and partner hubs (DIFC Dubai, Singapore VCC, Israel’s BIRD Foundation model, which has generated \$10B+ in indirect revenues from over 1,000 projects). The GCIZ becomes a node in a global trade finance network, not an isolated cluster.

Geopolitical intelligence. AI-mediated analysis of trade restrictions, sovereignty movements,

and regulatory divergence. Cambridge’s independent AI capability — community-owned models, no dependency on US or Chinese providers — makes this intelligence service credible to hubs concerned about technology sovereignty.

The GCIZ scales by making global trade investable. Its FinTech and InsurTech infrastructure produces tools and data products that financial institutions, trade finance providers, rating agencies, and sovereign wealth funds will pay for globally.

9 Intervention Network Analysis

10 Proposals by Layer

We may decompose each layer into a collection of intervention proposals following our frameworks towards scaling the ecosystem. We may also map out how interventions relate in a detailed manner, treating each as a central nodes, and the dependencies and outputs as star nodes, creating a node star graph for each intervention. With this, where outputs serve as inputs for other interventions, efforts become shared capabilities, from which we compose a full map of dependencies, generating a metaecosystem of ecosystem scaling activities. This is represented for each intervention as a node-star listing, and shown in a companion document.

10.1 Layer 1: Ecosystem Sensing

Automates: Knowing what exists, who needs what, who offers what — continuously.

10.1.1 1a. Live Venture Registry (Stages 1–2).

Ingest Companies House data, patent filings, funding announcements, hiring signals, and publications. Registration becomes a side-effect of activity (the Google Scholar model) rather than a separate administrative process. The external model is Israel’s IVC Research Center, which maintains comprehensive real-time tracking of every Israeli tech company, investor, and deal — tracking over 8,000 companies and 1,500 investors — providing the data layer that enables Israel’s ecosystem to present itself as a coherent network to global LPs. Cambridge lacks an equivalent. Note: even the registry’s identity layer requires financial infrastructure (KYC rails, verified credential systems).

Node-Star: 9 outbound connections (1b, 1c, 1d, 2a, 3a, 3c, 3e, 4c, 6a); 3 inbound feedback loops (3e, 4a, 4c). Key stakeholders: Cambridge Enterprise, Innovate Cambridge, Dealroom.

10.1.2 1b. Needs-Offers Matching Engine (Stages 2–3).

Every venture has needs (capital, talent, regulatory guidance) and offers (tech, expertise, data). Infer from signals, maintain a live index, and replace manual brokering. External

models: Singapore’s SGInnovate talent matching system and Enterprise Singapore’s GoBusiness platform.

Node-Star: In: Registry (1a), investor mandates, talent pool data. Out: matched pairs, gap analysis, demand signals. Feeds: 1c, 3a, 4a. Receives: 1a, 6b.

10.1.3 1c. Network Health Diagnostics (Stages 3–4).

Centrality measures, structural hole detection, community clustering (8) — making ecosystem structure visible and queryable. External model: digital twin methodology (Siemens, GE, Rolls-Royce). Startup Genome’s Global Startup Ecosystem Report already ranks Cambridge first for knowledge; this builds the live computational twin that doesn’t exist anywhere. The diagnostic tool itself becomes an exportable product at Stage 5.

Node-Star: In: 1a, 1b, deal outcomes. Out: health dashboard, bottleneck identification, benchmarking data. Feeds: 4a, 4b, 6d.

10.1.4 1d. Sector Emergence Detection (Stage 4).

Patent co-filing, hiring clusters, and co-authorship signalling new domains — breaking deeptech and medtech lock-in and detecting emerging FinTech and InsurTech opportunities within Cambridge’s research base. External models: StartUs Insights’ Discovery Platform, CB Insights Mosaic Scores (15). Stage 4 — Cambridge generates its own detection models, trained on local data, calibrated to its unique research-to-venture pipeline. Exportable.

Node-Star: In: 1a, patent/hiring/publication data. Out: sector alerts, diversification opportunities, investment signals. Feeds: 3b, 4b.

10.1.5 Scaling value.

- Compression: months of manual venture mapping become automatic and continuous.
- Amplification: pattern detection beyond human capacity; emergent sectors identified before they are visible.
- Resilience: knowledge decentralised — no single broker holds the ecosystem map.
- Revenue: data products licensed to other hubs as the sensing module.

10.2 Layer 2: Operational Automation

Automates: Governance, contracting, procurement, financial administration — the LegalTech and RegTech layer.

10.2.1 2a. Growth-Stage Governance Routing (Stages 1–3).

Venture profile in → prioritised governance roadmap out, across regulatory compliance, IP, data, financial, and institutional governance, optimised for ROI at the current stage. External models: RegTech sector (ComplyAdvantage, Ascent, Cube) — NLP parsing regulatory text mapped to business operations; Estonia’s X-Road — interoperability for verified information exchange without centralised control; Dubai’s DIFC Innovation Hub — regulatory sandbox with structured governance navigation. Naturalised for UK/Cambridge regulatory context, especially SEIS/EIS, university IP frameworks, and FCA cooperative rules.

Node-Star: In: 1a, regulatory corpora. Out: compliance roadmaps, pre-certification status, governance training data. Feeds: 2b, 3a, 3f. Receives: 1a, 4b, 5, 6b.

10.2.2 2b. Contract Generation and Procurement Compression (Stages 2–3).

AI-drafted NDAs, IP assignments, and investment terms calibrated to deal type. Automated vendor assessment and pre-certification of ecosystem ventures against corporate requirements. Target: 6–18 months → 6–18 weeks. External models: Accord Project (Linux Foundation) — open-source smart legal contract templates; Y Combinator’s SAFE and 500 Global’s KISS. Cambridge-specific instruments needed for university spinout deals, cooperative structures, and cross-hub agreements. Contract automation inherently involves payment terms, escrow conditions, and milestone-triggered disbursements — the operational layer naturally requires financial infrastructure. Key stakeholders: Cambridge law firms (HCR Hewitsons, Mills & Reeve), university legal offices, Accord Project community.

Node-Star: In: 2a, templates, payment infrastructure. Out: standardised instruments, pre-certified ventures. Feeds: 3a, 4a.

10.2.3 2c. Financial Operations Automation (Stage 2).

Invoicing, grant reporting, cap table management, SEIS/EIS claims, and R&D tax credits. External models: Carta — cap table management and private share administration serving 40,000+ companies; AngelList Stack — full back-office for venture funds. Stage 2 naturalisation: adapt US-origin tools for the UK tax regime, SEIS/EIS, and Cambridge’s SME profile.

Node-Star: In: cap table data, SEIS/EIS parameters. Out: clean financial records, optimised claims, freed founder time, financial performance data. Feeds: 3b, 5.

10.2.4 Scaling value.

- Compression: legal costs and procurement cycles cut and compressed for early ventures.
- Amplification: small deals become economical; corporate innovation opens to previously excluded ventures.
- Resilience: no single advisor or law firm is critical.

- Revenue: compliance-as-a-service, contract automation licensing, financial operations SaaS — all exportable service-economy products.

10.3 Layer 3: Intelligent Capital Coordination

Automates: Matching capital to opportunities, structuring vehicles, creating instruments that retain ecosystem value.

10.3.1 3a. Proactive Deal Flow Matching (Stages 1–3).

Continuous matching of investor mandates against the live registry, with matches surfaced proactively. External models: Israel’s Start-Up Nation Central — a searchable platform connecting global investors with Israeli startups; OurCrowd — combining venture investing with a technology platform. Highest inbound connectivity in the capital layer.

Node-Star: In: 1a, 1b, 2a, 2b, 3c, 4c (6 inbound). Out: deal pipeline, capital deployment velocity data. Feeds: 3b, 3d, 5, 6a.

10.3.2 3b. Dynamic Fund Structuring and Tax Arbitrage Optimisation (Stages 2–4).

Model optimal vehicles given SEIS/EIS parameters, investor profiles, and deal flow, adapting as incentive structures change. Cambridge’s ecosystem size is fundamentally linked to the magnitude of tax incentives — modelling this relationship quantitatively and optimising represents a novel contribution. External model: AngelList Rolling Funds and Syndicate structures (closest US analogy). Cambridge should build the UK-specific equivalent, incorporating SEIS/EIS optimisation. Stage 4 — Cambridge-native technology; the tax-arbitrage-to-ecosystem-size model is exportable to any jurisdiction with innovation tax incentives.

Node-Star: In: SEIS/EIS parameters, 3a, 1d, 2c. Out: optimised fund structures, exportable structuring tools. Feeds: 3f, 5.

10.3.3 3c. IP Valuation Engine (Stages 3–5).

Train on longitudinal Cambridge venture data, patent citation networks, academic impact metrics, team composition data, and comparable transactions. Produce defensible valuations for deep-tech IP that rating agencies currently cannot. External models: PatentSBERTa (transformer-based patent landscaping), Ocean Protocol data asset valuation concepts. No existing hub has solved this. Stage 5 — export to investors, TTOs, rating agencies, and pension funds globally. This is the single most globally exportable product the framework can produce.

Node-Star: In: 1a longitudinal data, citation networks, impact metrics. Out: defensible valuations, rating methodology, investable asset class. Feeds: 3a, 3e, 5. Receives: 6b.

10.3.4 3d. Value Retention Instruments (Stages 3–4).

Financial structures maintaining ecosystem stakeholdership through exits — convertible preferences, trailing royalties, ecosystem equity pools. External models: Novo Nordisk Foundation (commercial entity and foundation intertwined for perpetual stakeholdership); Israel’s Yozma programme (government co-investment model retaining national benefit from VC-backed exits). Addresses the £143 billion annual value leakage (31) by restructuring international linkages from extractive to bilateral.

Node-Star: In: 3a outcomes, cooperative legal frameworks. Out: stakeholdership through exits, recycled capital. Feeds: 5.

10.3.5 3e. Parametric Innovation Insurance (Stages 2–5).

Cover venture-specific risks — regulatory change, key person loss, IP challenge, market timing failure — with AI-driven assessment and automated payout based on measurable triggers rather than loss assessment. Extend to international port insurance and commodities-linked real estate insurance measured via insurtech data feeds (15; 16). Cambridge acts as both insured ecosystem and insurer of others. External models: Arbol (climate parametric insurance using blockchain for transparent trigger verification), Blink Parametric (travel/cyber), Corgi Insurance (\$108M funding, AI-native full-stack carrier for startups), Lloyd’s of London origin (merchants sharing risk, formalised). Growth: import parametric concepts (Stage 2); integrate with ecosystem sensing layer (Stage 3); develop Cambridge-native innovation insurance products (Stage 4); export as an insurtech product suite (Stage 5).

Node-Star: In: 1a sensing data, risk triggers, commodity price feeds. Out: risk coverage, automated payouts, premium income stream, insurance product suite for export. Feeds: 1a (feedback loop), 5.

10.3.6 3f. Cross-Hub Co-Investment Vehicles (Stages 1, 5).

Vehicles allowing Cambridge, London, Dubai, and Boston investors to co-invest with automated cross-jurisdictional compliance. External models: Dubai’s DIFC (common-law jurisdiction enabling structured cross-border investment); Singapore’s Variable Capital Company framework (flexible fund domiciliation for international investors). Stage 1 — establish bilateral links with DIFC, Singapore VCC, and Boston VC networks. Stage 5 — export the cross-hub vehicle architecture as a standard.

Node-Star: In: 4a agreements, 2a jurisdictional data, 3b structures. Out: bilateral capital flows, hub connectivity. Feeds: 4a, 5.

10.3.7 Scaling value.

- Compression: deal matching in days not months.

- Amplification: creates an investable asset class from previously unrated IP; enables parametric insurance for innovation risks that were previously uninsurable.
- Resilience: market-based valuation rather than expert-dependent; pooled risk across portfolio.
- Revenue: transaction fees, IP valuation tool sold globally, insurance premium income, fund structuring advisory.

10.4 Layer 4: Ecosystem-Level Dealmaking (Human-Led, AI-Supported)

What humans do once Layers 1–3 run.

The strategic-operational flip in practice: AI at functional levels (sensing, operations, capital) handles volume coordination, which is precisely what enables human dealmakers to operate at the ecosystem level rather than the transaction level. A dealmaker who no longer spends time on introductions, compliance checks, and contract drafting can spend that time negotiating a Cambridge–Dubai regulatory sandbox or restructuring a sovereign wealth fund relationship. At the company scale, AI handles operational dealmaking (matching, due diligence, term sheet generation); at the ecosystem scale, humans handle strategic dealmaking (partnership design, institutional negotiation, policy shaping). The two levels are complementary and interdependent.

10.4.1 4a. Hub-to-Hub Strategic Partnerships (Stage 5).

- Cambridge–London shared deal flow and co-investment;
- Cambridge–Dubai joint regulatory sandbox;
- Cambridge–Boston talent mobility and shared accelerators;
- Cambridge–Singapore sovereign capital access;
- Cambridge–Tel Aviv bilateral R&D.

Each partnership is formalised through co-investment vehicles from Layer 3 and governed by nested Ostrom principles. External model: Israel’s bilateral R&D agreements — BIRD Foundation (with US), SIIRD (Singapore), KORIL (Korea) — structured inter-hub partnerships with co-funding, shared IP, and defined benefit-sharing. BIRD has funded over 1,000 projects since 1977, generating \$10B+ in indirect revenues, converting Israel’s ecosystem from a local cluster into a globally networked asset class. Cambridge has no equivalent bilateral innovation infrastructure. AI contribution: Layers 1–3 provide the intelligence base — sensing identifies complementary capabilities between hubs, operations assesses regulatory compatibility, capital models co-investment structures. Human dealmakers use this intelligence to design and negotiate. Key stakeholders: GCI leadership, Innovate Cambridge, DIFC, EDB Singapore, BIRD Foundation.

Node-Star: In: 1b, 1c, 2b, 3f, 6e. Out: bilateral hub agreements, shared sensing data, co-investment routing, expanded ecosystem scope. Feeds: 1a (major feedback loop), 3f.

10.4.2 4b. Institutional and Policy Campaigns (Stage 5).

Lobbying for expanded SEIS/EIS; negotiating new academic pathways (PhD by Innovation) through university governance, requiring engagement with the Pro-Vice-Chancellor for Innovation, the Board of Graduate Studies, and progressive colleges (Hughes Hall, Lucy Cavendish); engaging ARIA (10-year roadmap: 5 years experimentation, 5 years commercialisation), Innovate UK, and British Patient Capital. Building fluid distribution chains so institutional rejection becomes redirection rather than termination. AI contribution: stakeholder mapping, incentive modelling, scenario analysis, calibrated communication drafting. Key stakeholders: Cambridge Enterprise, ARIA, Innovate UK, British Patient Capital, Cambridge Ahead.

Node-Star: In: 1c, 1d. Out: policy changes, new academic pathways, institutional funding, regulatory sandbox access. Feeds: 2a, 3b.

10.4.3 4c. Anchor Tenant Recruitment (Stages 1, 5).

Persuading corporates, sovereign wealth funds (Temasek, Mubadala, PIF — all with innovation ecosystem mandates), and international research organisations to establish Cambridge presence. External models: Abu Dhabi's Hub71 (incentive packages attracting 300+ startups from 50+ countries since 2019); Singapore's EDB (anchor-tenant strategy contributing to Asia's #1 startup ecosystem ranking). AI contribution: the sensing layer identifies which external organisations have strategic needs aligning with Cambridge capabilities; the capital layer models the value proposition in their terms; the operations layer pre-certifies Cambridge ventures against the anchor tenant's procurement requirements. Key stakeholders: Cambridge Ahead, CIC, DIT trade representatives, target SWF innovation arms.

Node-Star: In: 1a, 2b. Out: corporate and SWF presence in Cambridge, anchor investment, market access. Feeds: 1a (feedback loop), 3a, 5.

10.4.4 4d. Geopolitical and Trade Intelligence (Stages 4–5).

AI-mediated analysis of geopolitical dynamics affecting innovation flows — trade restrictions, sanctions, regulatory divergence, technology sovereignty movements. Cambridge's independent AI capability (not dependent on US/Chinese infrastructure) becomes a strategic asset for navigating complexity that changes faster than human analysts can track. No direct equivalent exists at ecosystem level; corporate geopolitical risk platforms (Dataminr, Predata) serve individual firms. Cambridge can build the ecosystem-level equivalent, where the sensing layer's data feeds include geopolitical signals alongside venture and market data. Exportable service for partner hubs navigating US-China decoupling. Key stakeholders: Centre for Science and Policy (CSaP), Cambridge AI research groups, IfM supply chain research.

Node-Star: In: 1a, geopolitical and trade data feeds. Out: risk reports, dealmaking intelligence, exportable intelligence service. Feeds: 4a.

10.4.5 Scaling value.

- Compression: AI-prepared intelligence for every meeting; dealmakers arrive armed rather than starting from scratch.
- Amplification: ecosystem-to-ecosystem deals, sovereign wealth fund engagement, global trade navigation become possible.
- Resilience: multiple hub relationships, no single partnership critical.
- Revenue: co-investment vehicle fees, advisory services, trade intelligence subscription.

10.5 Layer 5: Financial Self-Sustainability

Ensures: Framework pays for itself and generates surplus through revenue mechanisms embedded in coordination activity.

The GCIZ's financial sustainability is not a separate concern bolted onto the coordination infrastructure — it is the coordination infrastructure expressed in financial terms. Each revenue mechanism emerges directly from an operational layer, which means revenue scales with usage rather than requiring independent sales effort. The seven mechanisms below form a staged portfolio: early mechanisms (membership, compliance) generate cash flow while the infrastructure is being built; later mechanisms (data products, IP licensing, SaaS) represent the high-margin returns on accumulated capability. External models informing the overall architecture: platform economics (43) — two-sided markets where value accrues to the orchestrator of network effects; embedded finance — financial services delivered at point of need within non-financial workflows, eliminating friction between coordination and payment; Lloyd's of London — merchants sharing risk, formalised into a market that became a global standard.

10.5.1 5a. Membership Tiers — Community and Full, with Mutual Stake (Stages 1–5).

Tiered membership providing differentiated access to ecosystem services: community tier offers registry visibility, event access, and basic matching; full tier adds compliance routing, capital coordination, insurance eligibility, and governance participation. Mutual stake means members hold cooperative equity, aligning incentives and reducing churn. External models: London Stock Exchange Group's membership structure (tiered access to market infrastructure); WeWork Labs (tiered startup services, though without cooperative ownership); the Co-operative Group (mutual ownership at commercial scale, distributing surplus to members). The cooperative element is not cosmetic — it is structurally necessary for data sharing, since ventures will contribute operational data to the sensing layer only if they retain ownership stakes in the data asset that results.

Node-Star: In: 1a (registry data for onboarding). Out: recurring revenue, member data enriching 1a, governance participation rights. Feeds: all layers (membership revenue funds infrastructure), 1a (feedback loop through onboarding data).

10.5.2 5b. Compliance-as-a-Service (Stages 2–5).

Productising the governance routing engine (2a) and contract generation capability (2b) as a subscription service for ecosystem members and external clients. Regulatory mapping, automated SEIS/EIS certification, FCA compliance checking, and IP governance audits delivered as on-demand services. External models: ComplyAdvantage (AML/KYC compliance-as-a-service, valued at over \$1B); Avalara (tax compliance automation); Stripe Atlas (company formation and compliance bundled into a platform). Cambridge’s naturalised regulatory intelligence — built for UK tax, university IP, and cooperative structures — is not replicable by generic RegTech providers. Stage 2: internal service for GCIZ members. Stage 3: external sales to Cambridge-adjacent firms. Stage 5: licensed to partner hubs adapting the compliance engine to their jurisdictions.

Node-Star: In: 2a (governance routing output), 2b (contract templates), regulatory corpora. Out: subscription revenue, compliance certification data, regulatory intelligence. Feeds: 2a (feedback through usage data improving routing accuracy), 5g (compliance module as component of ecosystem OS).

10.5.3 5c. Transaction Fees on Capital Coordination (Stages 3–5).

A percentage fee on capital deployed through the ecosystem’s matching and structuring infrastructure — deal flow matching (3a), fund structuring (3b), and cross-hub co-investment vehicles (3f). The fee structure follows a declining marginal rate: higher for smaller deals (where the platform’s value-add relative to manual alternatives is greatest) and lower for larger transactions (where volume justifies thinner margins). External models: AngelList (carry and platform fees on syndicated deals); OurCrowd (management and success fees on a venture investing platform); Stripe (percentage-based transaction fees embedded invisibly in the workflow). The critical design principle is that fees must be low enough that the platform remains cheaper than the manual alternative at every deal size.

Node-Star: In: 3a (deal flow), 3b (fund structures), 3f (cross-hub vehicles). Out: transaction revenue, deal volume data. Feeds: 3a (revenue enabling continued platform development), 6a (transaction data enriching training sets).

10.5.4 5d. Parametric Insurance Premium Income (Stages 3–5).

Revenue from premiums collected on parametric innovation insurance products (3e), where Cambridge acts as both insured ecosystem and insurer of partner hubs. Premium income is the direct revenue stream; the actuarial data generated is the indirect asset. External models: Lloyd’s of London (syndicates earning premium income while building underwriting expertise);

Swiss Re (reinsurance as a data-driven business); Lemonade (AI-native insurance with low marginal cost per policy). Growth: Stage 3 — internal ecosystem policies generating modest premium pools; Stage 4 — Cambridge-developed innovation risk models attracting external insurers as co-underwriters; Stage 5 — export of parametric insurance products to partner hubs, with Cambridge retaining actuarial IP. The premium income stream is particularly important because it is counter-cyclical: during downturns, when other revenue streams shrink, insurance demand increases.

Node-Star: In: 3e (parametric insurance infrastructure), 1a (risk trigger data). Out: premium income, actuarial datasets, counter-cyclical revenue buffer. Feeds: 3e (premium pool enabling broader coverage), 6a (actuarial data for AI training).

10.5.5 5e. Data Products — Anonymised Ecosystem Intelligence (Stages 3–5).

Anonymised, aggregated datasets and intelligence reports derived from the ecosystem’s operational data: venture trajectory patterns, sector emergence signals, deal outcome distributions, compliance pathway benchmarks, and risk event correlations. External models: CB Insights (ecosystem intelligence products sold to corporates and investors); PitchBook (private market data as a subscription business); Bloomberg Terminal (financial data infrastructure commanding premium pricing through comprehensiveness). Cambridge’s data is structurally different from these — it includes internal operational dynamics (not just public filings and press releases) because ventures contribute data as cooperative members. Stage 3: initial reports for GCIZ members. Stage 4: subscription products for external investors and corporates. Stage 5: real-time data feeds and API access.

Node-Star: In: 1a (sensing data), all layer operational outputs. Out: subscription revenue, ecosystem benchmarking data, investor intelligence. Feeds: 6a (curated datasets for AI training), 4c (data products as anchor tenant recruitment tool).

10.5.6 5f. IP Valuation Tool Licensing (Stages 4–5).

The IP valuation engine (3c) is the single most globally exportable product the framework produces. Licensing it to investors, technology transfer offices, rating agencies, insurers, and pension funds as a standalone valuation service. External models: Moody’s and S&P (credit rating methodologies licensed globally); FICO (scoring model built on longitudinal data, now underpinning a multi-trillion-dollar lending industry); Derwent Innovation (patent analytics platform used by IP-intensive industries). The tool’s defensibility comes from its training data: no other hub has instrumented infrastructure generating longitudinal venture-to-exit data with full context (team, technology, market, regulatory, governance). Stage 4: beta licensing to Cambridge investors and TTOs. Stage 5: global licensing with jurisdiction-specific calibration modules.

Node-Star: In: 3c (valuation engine), longitudinal ecosystem data. Out: licensing revenue, valuation methodology as industry standard, deep-tech investability. Feeds: 3c (licensing revenue

funding continued model improvement), 4a (valuation tool as partnership asset for hub-to-hub deals).

10.5.7 5g. SaaS Licensing of Ecosystem OS (Stage 5).

The culminating revenue mechanism: licensing the entire integrated platform — sensing, operations, capital coordination, governance, and AI layers — as a configurable software-as-a-service product for other innovation hubs globally. External models: Salesforce (CRM as the template for enterprise SaaS at scale); Shopify (commerce OS enabling others to build on shared infrastructure); Palantir Foundry (complex operational platform sold to governments and institutions). The ecosystem OS is not a single application but a modular suite: hubs can license individual layers (sensing only, or compliance only) or the full stack. Each deployment generates data that feeds back into Cambridge’s AI training pipeline, creating a network effect where more licensees produce better models for everyone. Stage 5 assumes at least 2 partner hubs have been operating on the platform for 12+ months, providing evidence of cross-hub portability.

Node-Star: In: all layers (integrated platform output), 4a (hub partnerships as distribution channel). Out: SaaS revenue, global deployment data, network effects across hubs. Feeds: 6a (multi-hub data enriching training sets), 4a (SaaS clients become formal hub partners).

10.5.8 Scaling value.

- Compression: revenue collection embedded in coordination workflows, approaching zero marginal cost per additional transaction as the platform matures.
- Amplification: seven revenue streams diversify income across stages, sectors, and geographies — no single mechanism must succeed for the system to be financially viable; data products and IP licensing create high-margin revenue uncorrelated with local deal volume.
- Resilience: cooperative ownership prevents acquisition or extraction of the revenue infrastructure; counter-cyclical insurance income buffers downturns; SaaS licensing distributes operational risk across multiple hubs.
- Revenue: the portfolio is designed so that early mechanisms (membership, compliance) are cash-flow positive by Stage 2, mid-stage mechanisms (transaction fees, insurance premiums, data products) reach sustainability by Stage 3–4, and high-margin mechanisms (IP licensing, SaaS) generate surplus from Stage 5. The GCIZ pays for itself because the fintech layer is not separate from the coordination layer — it is the coordination layer expressed in financial terms.

10.6 Layer 6: AI Ecosystem Generation

Produces: AI tools, models, and services that serve the framework and become independently valuable.

10.6.1 6a. Training Data Creation (Stage 3 onward).

Each layer generates labelled datasets that currently don't exist: venture trajectories, governance mapping corpora, deal outcomes, negotiation patterns, insurance events, and trade flow data. No other hub generates these datasets because no other hub has instrumented infrastructure. The dataset comprises data that accumulates from day one of ecosystem OS operation.

Node-Star: In: all layer operational data. Out: labelled datasets, longitudinal ecosystem corpus, unique data asset. Feeds: 6b, 6d, 6e.

10.6.2 6b. Demand Pull on Cambridge AI Research (Stages 3–4).

Unsolved framework problems become research opportunities pulling Cambridge AI groups toward applied service-economy problems. Procurement compression needs contract NLP; IP valuation needs multimodal models; hub matching needs cross-cultural AI; supply chain risk needs specialised models. Each connects the ecosystem OS directly to Cambridge's research base.

Node-Star: In: unsolved problems from all layers (contract NLP from 2b, valuation modelling from 3c, cross-cultural matching from 4a, supply chain risk from 4d); training data (from 6a); Cambridge AI research capacity (CSIT, Machine Learning Group, IfM). Out: improved AI models, published research, enhanced layer performance, talent attraction pipeline. Feeds: 1b (improved matching models), 2a (improved regulatory NLP), 3c (improved valuation models), 4d (improved geopolitical analysis). Receives: 6a (training sets providing the data on which applied research is conducted).

10.6.3 6c. Controlled Deployment Environments (Stage 4).

Bounded autonomy architectures: operational limits, human escalation, audit trails. The ecosystem OS provides a production testbed for agentic AI with real stakes but structured governance. PwC's 2025 Responsible AI survey finds that only around 28% of US business leaders place their organisations at the most advanced ("strategic") stage of AI governance maturity, with the remainder still embedding, training, or building foundational policies (46); Cambridge's Ostrom-grounded cooperative governance provides a structural advantage in this critical gap. Cambridge's AI safety research community gains a living laboratory.

Node-Star: In: integrated layer infrastructure (Stages 1–3 outputs); AI models (from 6b); Ostrom governance rules; audit and escalation protocols. Out: validated deployment architectures, structured feedback data, governance-tested agentic systems, living laboratory for AI safety research. Feeds: 6d (validated models entering ecosystem control), 6e (deployment performance data enriching prediction training), 5 (governance-as-a-service credibility). Receives: 6a (training data for test deployments), 6b (candidate models requiring production validation).

10.6.4 6d. Ecosystem Control Mechanism (Stage 4).

Accumulated infrastructure from Stages 1–3 requires its own control mechanism — a bespoke AI model and technology suite for managing the integrated ecosystem in real time. This is the Cambridge-native product: an AI system trained specifically on the dynamics of innovation ecosystems, capable of coordinating sensing, operations, capital, and governance layers simultaneously.

Node-Star: In: 1c, 6a, 6b. Out: coordinated ecosystem management, automated layer orchestration, Cambridge-native AI product. Feeds: 1a, 5.

10.6.5 6e. The Innovation Prediction Model — Moonshot Output (Stages 4–5).

If Cambridge runs the ecosystem OS for enough years, it will generate a dataset comprising: venture formation events with full context (founder profiles, technology characteristics, market conditions, regulatory environment); compliance pathway choices and outcomes; capital matching events and subsequent performance; team composition changes and effects; cross-hub collaboration patterns and results; and governance decisions and ecosystem health impacts. This dataset does not exist anywhere because no ecosystem has the instrumented infrastructure to generate it. The model trained on this data — "given these conditions, what is the probability distribution of outcomes?" — would be a foundational contribution to the science of innovation, with a similar type of contribution to what genomic databases did for personalised medicine (38). The closest existing analogue is CB Insights' Mosaic Score (15), but that is a scoring system on public data, not a foundational model on internal ecosystem dynamics. Credit scoring models (FICO) were trained on decades of longitudinal financial data and now underpin a multi-trillion-dollar lending industry. The ecosystem prediction model likely requires 3–5 years of instrumented operation across 500+ ventures to produce initial utility, reaching full capability at 5–10 years. The data accumulates from day one. This is the moonshot justification for the entire infrastructure investment: even if individual layers generate modest standalone revenue, the accumulated data asset justifies the enterprise.

Node-Star: In: 5+ years longitudinal data from 6a, all layer outputs. Out: outcome probability distributions, portfolio construction tool, policy design tool, globally exportable product. Feeds: 5, 4a.

10.6.6 6f. Open-Source Release Strategy (Stage 5).

Strategic open-sourcing of framework components. The Hugging Face model: open-source becomes the default; commercial value flows through the enterprise layer. The CERN model: research infrastructure producing spin-offs (the World Wide Web). Cambridge positions itself as the origin of innovation ecosystem AI.

Node-Star: In: mature framework components (from all layers); ecosystem control models (from 6d); open-source packaging and documentation; commercial licensing terms. Out: publicly

available ecosystem AI tools, developer community formation, Cambridge brand as origin of innovation ecosystem AI, enterprise licensing revenue. Feeds: 4a (open-source adoption as hub partnership entry point), 5 (enterprise-layer licensing revenue), 1a (community contributions expanding registry coverage). Receives: 6d (production-validated models ready for release), 5g (SaaS architecture defining the enterprise layer built on top of open-source components).

10.6.7 Scaling value.

- Compression: AI improves every layer’s efficiency with each cycle, approaching near-zero marginal cost per additional coordination task (54).
- Amplification: innovation prediction models and domain-specific agentic architectures that no other hub can produce.
- Resilience: open-source ensures continuity beyond any single institution; community-owned models resist platform lock-in.
- Revenue: AI tools sold to financial institutions, TTOs, rating agencies, trade finance providers, and policymakers; SaaS licensing; research partnerships.

11 Conclusions

The GCIZ exports not platforms or hardware, but ecosystem infrastructure, risk products, and trade intelligence — the service layer that sits between technology for global players and the rest of the world’s innovation economies.

11.1 Alternative Approaches and Cambridge’s Unique Conditions

Other hubs have addressed similar coordination challenges through different models, each shaped by local conditions. Estonia’s e-Residency and X-Road is state-led, requiring unified governance that Cambridge — spanning university, private sector, and multiple levels of government with no single authority — cannot replicate. Singapore’s Smart Nation and EDB strategy requires central planning, massive state capital reserves, and a small geographic footprint. Israel’s Start-Up Nation Central and BIRD Foundation emerged under existential security pressures that created a culture of compulsory cooperation and risk tolerance; Cambridge’s culture is collegiate and voluntary. Silicon Valley’s market-led model depends on massive capital concentration, a winner-take-all culture, and decades of network accumulation; Cambridge’s SME structure, academic culture, and UK capital environment are fundamentally different.

Cambridge’s unique conditions — world-leading research density in a compact geography, a cooperative rather than competitive culture, strong academic institutions, UK legal infrastructure for cooperatives and mutuals, a fintech research base without a fintech company base, and greenfield digital infrastructure — require a model that draws from each of these precedents but is identical to none.

11.2 Relationship to Prior Global Cambridge Models

This paper closes a trilogy. *The GCI Strategy Plan* (?) defined *what Cambridge needs to build* — a government-owned, GCI-operated investment zone combining tax-advantaged policy, curated membership, shared services, and public dashboards, designed to counter the structural extraction documented there. *Systematising Startup Creation* (44) defined *how ventures are produced within that zone* — the Cambridge Startup Production Model, a five-layer system of Training, Capital, Production, Currency, and Culture designed to convert Cambridge’s research base into a generator of entrepreneurship, and to do so without relying on the ad-hoc individual initiative that has historically carried the ecosystem. This paper defines *how the zone itself scales as a coordinated economic system* — not one firm at a time, but as a platform that coordinates ventures, capital, and cross-hub relationships through agentic infrastructure and a fintech-native substrate.

The three papers form a progression from container, to contents, to coordination. The first asks what institutional, legal, and fiscal shell can keep value in Cambridge. The second asks what to put inside that shell — what kinds of ventures, produced how, by whom — given that the shell on its own does not create companies. The third asks how what is inside the shell grows, connects outward, and compounds over time without coordination costs rising in proportion to volume. Each paper extends the last: the Production Model builds on the zone, and the ecosystem operating system builds on both.

Continuity across the three is strongest on the question of retention. The first paper framed retention as a policy-and-incentives problem and argued that without a zone, ventures will continue to optimise toward the jurisdictions that offer the cleanest late-stage capital. The second framed retention as a talent-and-culture problem and argued that without deliberate production of founders with both research depth and commercial ambition, the incentive structure has too little to act on. This paper frames retention as a coordination-infrastructure problem and argues that even with the right tax regime and a pipeline of commercially ambitious founders, Cambridge cannot scale the ventures it produces if its coordination layer remains emails, spreadsheets, and personal introductions across its web of active companies, research parks, and innovation-driven firms. The three interventions are complements, not alternatives: each addresses a distinct axis along which value leaks, and removing any one of them would leave the others materially less effective.

In contrast to the previous papers, which treat human and physical infrastructure as primary intervention targets, this paper focuses on tying together the hubs of human and physical infrastructure that already exist — or that are proposed in the earlier reports — and creating digital autonomy over their coordination. *Systematising Startup Creation* addressed talent extensively through its Training Layer; those proposals remain central and are not repeated here. Likewise, the physical constraints and growth opportunities discussed in the earlier work — including housing, laboratory and office space, and corridor-scale developments such as the Oxford–Cambridge Growth Corridor and Cambridge Innovation Hub (35) — are treated as

complementary inputs rather than objects of redesign. The ecosystem OS lowers coordination costs; it does not lower rent. What this paper adds is the scaling operating layer that connects these prior interventions into a single, agentic system.

What this paper uniquely adds, and what the trilogy as a whole depends on, is the question of *who builds and operates the scaling infrastructure itself*. This operating layer is not neutral plumbing but the principal mechanism by which Cambridge moves from being a research cluster that feeds other economies to being an economy in its own right. The zone retains; the production model produces; the operating system scales. Together they constitute a full pathway from invention, to venture, to ecosystem.

11.3 Frameworks developed

This paper developed a collection of conceptual and operational frameworks that can be used as a blueprint for designing a scalable innovation ecosystem within Cambridge:

1. **Triple Definition of Ecosystem Scaling** — Scaling decomposed into three simultaneous processes: compression (reducing unit coordination cost toward zero), amplification (enabling qualitatively new operations impossible at craft scale), and resilience (eliminating single points of failure).
2. **Five-Stage Growth Pathway** — Establish → Naturalise → Integrate → Generate → Export, with measurable thresholds governing each stage transition. Stages move from importing external tools, to adapting them locally, to achieving system coherence, to producing Cambridge-native AI, to global licensing.
3. **Six Infrastructure Layers** — A layered architecture for ecosystem coordination: Layer 1 (Sensing), Layer 2 (Operations), Layer 3 (Capital), Layer 4 (Dealmaking), Layer 5 (Financial Sustainability), Layer 6 (AI Generation). Lower layers are automated; upper layers are human-led, AI-supported.
4. **The Agentic Innovation Ecosystem** — A new organisational category defined as a geographically anchored but globally connected network of autonomous entities whose routine coordination is performed by purpose-built AI agents within cooperative governance, sustained by embedded financial infrastructure, and producing AI as both byproduct and primary export. Distinguished from the agentic organisation (single firm) by having no single governing authority and participants who can exit at will.
5. **FinTech and InsurTech as Organising Substrate** — Financial technology repositioned from a sector to the medium through which every ecosystem relationship is formalised, decomposed into four sub-layers: social infrastructure (membership, mutual stakes), business infrastructure (deal flow, co-investment), monetary infrastructure (fees, premiums, licensing), and technical infrastructure (payment rails, KYC/AML, smart contracts).
6. **The Data Flywheel as Elevation Mechanism** — Each layer's operations generate

data that trains AI that improves every layer. Functions simultaneously as the mechanism of progression through growth stages and the system’s primary export product, culminating in the Innovation Prediction Model.

7. **Nine Governing Principles** — Automate Volume/Liberate Judgment; Strategic-Operational Flip; Distribution Before Technology; Governance as Competitive Advantage; FinTech as Substrate; Multi-Hub Connectivity as Design Primitive; Portability as Deal-making Mechanism; Data Flywheel as Elevation Mechanism; Trust is Architectural.
8. **The Scaling Loop** — A closed feedback architecture connecting all six layers, where each layer’s output becomes the next layer’s input, compounding geometrically because each layer’s improvement enhances every other layer’s performance.
9. **Intervention Network Architecture** — 29 interventions organised as node-star graphs with explicit inbound dependencies and outbound connections, composing into a meta-ecosystem map of scaling activities. Each intervention is tagged to specific growth stages.
10. **Cooperative Governance Architecture** — commons principles mapped to specific ecosystem design features (defined boundaries → tiered membership; proportional equivalence → mutual stake tiers; nested enterprises → hub-within-hub governance; etc.), implemented under the UK Co-operative and Community Benefit Society Act 2014.

11.4 Key Observations

A collection of key observations can be derived from this paper summarised into the following:

1. **Scaling has a specific, measurable definition — and the existing literature doesn’t provide one.** Scaling is three simultaneous processes: compression (reducing unit coordination cost toward near-zero marginal cost), amplification (enabling qualitatively new operations impossible at craft scale), and resilience (eliminating single points of failure). This triple definition fills a gap in the entrepreneurial ecosystem, regional innovation system, and platform ecosystem literatures, none of which operationalise ecosystem-level scaling (37; 52; 30; 27). The five-stage growth pathway (Establish → Naturalise → Integrate → Generate → Export) is the mechanism by which this scaling is achieved, with each stage producing measurable advances on all three dimensions.
2. **The data flywheel is both facilitator and goal.** The framework generates data that trains AI that improves the framework — but the AI produced is also Cambridge’s primary export product. The infrastructure investment is therefore justified twice over: once as ecosystem scaling infrastructure, once as an AI development programme whose production environment is an innovation ecosystem. The innovation prediction model — trained on longitudinal ecosystem data no other hub possesses — is the moonshot output that justifies the entire enterprise.

3. **Fintech and insurance are the organising substrate, not a bolt-on.** Financial instruments formalise social connections (membership, mutual stakes), business connections (deal flow, co-investment), monetary connections (transaction fees, premiums), and technical connections (payment rails, smart contracts, identity systems). Every layer runs on this substrate. The framework pays for itself because the fintech layer is not separate from the coordination layer — it is the coordination layer expressed in financial terms.
4. **The agentic innovation ecosystem is a new organisational category.** Beyond the agentic firm, agentic mesh, and innovation hub concepts, the agentic innovation ecosystem is cooperative-governed, AI-coordinated, financially self-sustaining, and designed for a network with no single governing authority and participants who can exit at will.
5. **Cambridge has a greenfield advantage.** The three infrastructure obstacles — legacy integration, data discoverability, governance frameworks — explain why it is predicted that almost half of all agentic AI projects will be cancelled by 2027. Cambridge’s ecosystem OS is greenfield: no legacy systems, data architecture designed from scratch, governance built on community principles rather than retrofitted onto corporate hierarchy. This structural advantage enables leapfrogging established ecosystems locked into legacy coordination infrastructure.
6. **Geopolitical AI sovereignty is a selling point, not merely a risk mitigation.** The concentration of the AI stack among global players creates dependency for every other hub. Cambridge’s independent AI capability — developed through the ecosystem OS, owned by the cooperative, deployable without foreign strings — is attractive to any hub concerned about technology sovereignty.

11.5 Next Steps

Realising the framework now depends on coordinated execution across GCIZ stakeholders, government entities, and partner hubs, rather than further conceptual design. The actions and indicative timelines set out below are propositional, intended as a planning scaffold rather than fixed commitments. To transition the Global Cambridge Investment Zone (GCIZ) from a conceptual framework into a scalable, operational economic system, immediate practical implementations must run in parallel with advanced academic and technical refinements. The following outlines the roadmap for executing this vision and the ongoing research required to sustain its evolution.

11.5.1 Actions for GCIZ Stakeholders (Operators, Local Institutions, Investors)

GCIZ stakeholders should focus on building a minimal but complete “Stage-1 backbone” and then expanding via dependency-driven waves of interventions.

Establish an Implementation Board (proposed horizon: 0–3 months)

- Form a GCIZ Implementation Board drawn from Cambridge Enterprise, Innovate Cambridge, CIC, CCAF, major research institutes, and 1–2 early partner hubs.
- Give the board a mandate to move the ecosystem from Stage 1 to the threshold of Stage 2, with explicit responsibility for tracking the formal stage criteria already defined in the paper.
- Agree on decision-making norms, reporting rhythms, and how the board interfaces with existing governance bodies in Cambridge.

Define and Monitor Stage–Aligned KPIs

- For each initial intervention, specify a small set of metrics tied directly to the scaling definition (compression, amplification, resilience).
- Examples of target metrics (to be refined locally) include:
 - Compression: percentage reduction in time to find a relevant investor compared to today’s manual baseline.
 - Amplification: growth in the number of ventures and investors reachable through the platform versus the initial pilot group.
 - Resilience: diversity of data sources feeding the registry (number of institutions and contributors, not just volume).
- Review these metrics regularly (e.g. quarterly) and use them to decide when to deepen or pause specific interventions.

Build a Dependency–Based Roadmap and Delivery Waves

- Translate the intervention list into a visible implementation map that captures prerequisites and “unlocks”.
- Organise the work into approximately three waves, for example:
 - Wave 1 (Stage 1–2 focus): registry, basic needs–offers, governance routing, simple contract templates, initial deal–matching, and cooperative membership.
 - Wave 2 (Stage 3–4 focus): network diagnostics, sector emergence detection, parametric innovation insurance prototypes, cross–hub co–investment vehicles, ecosystem–level dealmaking, and early versions of the ecosystem control mechanism and innovation–prediction model.
- For each wave, GCIZ stakeholders should define “entry conditions” (what must be true to start) and “exit conditions” (which formal stage thresholds are satisfied).

Embed Stage Gates into Internal Funding and Governance

- Design GCIZ budgets so that funding for later-stage or higher-risk interventions is only released once specific stage criteria have been met (e.g. a 50% time/cost reduction for at least one coordination function before scaling certain Layer-3 tools).
- Encode these stage gates into cooperative or company governance documents wherever possible, so that moving from Stage 1 to Stage 2, or from pilot to full deployment, is a shared, rule-based decision rather than an informal judgment.

Run Structured Pilots and Public Post-Mortems Run tightly scoped pilots for each new intervention (e.g. registry + matching, governance routing, membership and contribution tracking), each with:

- A defined cohort, duration, and success measures.
- A clear plan for either graduating the intervention into production or cycling the learning back into further research and redesign.

Institutionalise Learning (Methods and Data)

- Create an internal “methods” function responsible for standardising data collection, defining common metrics, and ensuring that each intervention contributes usable data to the eventual AI models.
- This function should maintain the mapping between interventions, stages, and scaling metrics, and report annually on progress and bottlenecks.

11.5.2 Actions for Governmental Entities (Local, National, Regulatory)

Government actors—local authorities, national departments, regulators, and public funding bodies—shape the enabling environment. The dates and horizons mentioned below are indicative only; they are meant to suggest sequencing, not fixed deadlines.

Recognise GCIZ as a Strategic Testbed (proposed horizon: 0–12 months)

- Formally designate the GCIZ as a pilot environment for AI-enabled service infrastructure, FinTech/InsurTech innovation, and cross-hub coordination.
- Use existing mechanisms where possible (e.g. innovation zones, industrial strategy initiatives, AI assurance roadmaps) to reduce the need for bespoke legislation.

Provide Regulatory Sandboxes and Safe Harbours

- Expand or adapt regulatory sandbox programmes (for example in financial services and data governance) to explicitly cover GCIZ interventions such as parametric innovation insurance, cross-hub investment vehicles, and automated governance routing.

- Offer time-limited safe-harbour conditions for specific pilots, provided that risk-mitigation and audit mechanisms are in place.

Support Data and Infrastructure Access

- Facilitate access to relevant public datasets (corporate registries, innovation funding data, export statistics) under clear governance arrangements, so that the live venture registry and diagnostics tools can operate effectively.
- Where appropriate, co-invest in shared digital infrastructure (compute, secure storage, identity services) that will underpin the OS and its AI components.

Use GCIZ Metrics in Policy Evaluation

- Incorporate the GCIZ’s stage-based metrics (e.g. compression in coordination costs, resilience indicators, local vs. foreign ownership trajectories) into policy evaluation frameworks for innovation and industrial strategy.
- This enables government to treat the GCIZ not only as a beneficiary of policy but as a measurement instrument for what is working.

11.5.3 Actions for Partner Hubs (International and Domestic)

Partner hubs—both UK and international—are essential to realising the “agentic innovation ecosystem” as a genuinely networked concept. The sequence below is illustrative; individual hubs can adjust pace and depth based on their own priorities.

Define the Partnership Model and Scope (proposed horizon: 0–6 months for first pilots)

- Agree with GCIZ stakeholders whether a partner hub will initially engage as:
 - A data-sharing collaborator (limited registry interoperability and signals exchange).
 - A co-design partner for specific tools (e.g. cross-hub co-investment vehicles, joint parametric products).
 - A pathfinder for eventual licensing of the ecosystem OS (Stage-5 vision).
- Clarify expectations on governance, data rights, and the handling of cross-border regulatory issues.

Participate in Early Cross-Hub Interventions

- Join pilots of 3f. Cross-Hub Co-Investment Vehicles and 4a. Hub-to-Hub Strategic Partnerships once the GCIZ Stage-1 backbone is working.

- Contribute test cases, legal expertise, and local investor mandates to help stress-test interoperability.

Mirror Selected Interventions Locally

- Where feasible, adopt simplified versions of key GCIZ interventions (e.g. local venture registry schema, governance routing templates) so that data and workflows can be aligned across hubs.
- This creates the conditions for the later licensing or co-development of a shared ecosystem OS.

Co-Invest in the Research and AI Layers

- Collaborate on the design and training of the innovation prediction model and related AI components, especially by contributing anonymised longitudinal data where governance permits.
- Explore joint funding for Layer-6 work (training data creation, controlled deployment environments, ecosystem control mechanisms), with clear agreements on intellectual property and usage rights.

Prepare for Configurable OS Licensing (Stage-4/5 Horizon)

- As GCIZ approaches Stage 4-5, partner hubs should begin planning how a configurable version of the ecosystem OS would integrate with their own governance, financial, and regulatory structures.
- This includes assessing local cooperative or platform entities that could operate the OS, as well as potential revenue-sharing and risk-sharing models.

All dates and time horizons in this section are propositional and intended to illustrate a plausible sequence of work, not to prescribe binding deadlines. They can and should be adapted as GCIZ stakeholders, governmental entities, and partner hubs converge on their respective capacities, constraints, and priorities.

11.6 Next Research Steps

We close this with a brief proposal of further research that, upon being conducted, enriches and further enables the application of the concepts proposed in this paper: While the five-stage growth pathway provides a robust operational roadmap, advancing the concept of an "agentic innovation ecosystem" requires targeted academic research across several intersecting domains:

- **Parametric Innovation Finance:** Deeper actuarial and economic research is required to accurately model and price parametric insurance products tailored for high-growth startups. Future studies should focus on defining and quantifying complex risk triggers—such as

sudden regulatory shifts, supply chain blockages, or key-person loss—to ensure these instruments provide reliable, counter-cyclical revenue buffers.

- **Agentic Interoperability and Data Ontologies:** Because the ecosystem relies heavily on Layer 1 (Sensing) and Layer 2 (Operational Automation), computer science research must focus on standardizing data ontologies. Advancements in this area will allow federated AI agents to seamlessly interpret multi-jurisdictional legal texts, disparate patent filings, and complex financial data without relying on a centralized bottleneck.
- **Algorithmic Commons Governance:** Building upon Ostrom’s principles of common-pool resources, further socio-technical research should explore how to govern AI models within a mutual structure. This includes refining automated dispute resolution protocols, transparent contribution-tracking ledgers, and mechanisms for equitably distributing data dividends among tiered cooperative members.
- **Longitudinal Metrics for Ecosystem Scaling:** To empirically validate the paper’s proposed triple definition of scaling (compression, amplification, and resilience), researchers must design long-term longitudinal studies. By rigorously tracking the time and cost reductions of coordination events over several years, the ecosystem can refine the "innovation prediction model" and mathematically prove the efficacy of the data flywheel.

By treating the ecosystem itself as a continuous research and development environment, Cambridge can systematically scale its own venture production while exporting an empirically validated, AI-native operating system to peer hubs globally.

References

- [1] Cambridge can: Regional ai statement. Technical report, University of Cambridge, Mar. 2025. URL: <https://www.ai.cam.ac.uk/reports/cambridge-can-regional-ai-statement/>.
- [2] R. Adner. Ecosystem as structure: An actionable construct for strategy. *Journal of Management*, 43(1):39–58, 2017. Online November 2016. doi:10.1177/0149206316678451.
- [3] Atlantic.vc, Cambridge Innovation Capital, Dealroom.co, MITO Technology, Northern Gritstone, and Oxford Science Enterprises. European spinouts report 2025. Technical report, 11 2025. URL: <https://dealroom.co/reports/european-spinouts-report-2025>.
- [4] D. B. Audretsch, J. A. Cunningham, D. F. Kuratko, E. E. Lehmann, and M. Menter. Entrepreneurial ecosystems: economic, technological, and societal impacts. *The Journal of Technology Transfer*, 44(2):313–325, 2019. URL: <https://link.springer.com/article/10.1007/s10961-018-9690-4>, doi:10.1007/s10961-018-9690-4.
- [5] D. Barile, G. Secundo, and P. Del Vecchio. An artificial intelligence-based innovation ecosystem enabling open innovation and sustainable growth: Evidence from a case study. *Innovation: The European Journal of Social Science Research*, 2025. doi:10.1080/14479338.2025.2514468.
- [6] Beauhurst, Cambridge Enterprise, Innovate Cambridge, and Cambridge Innovation Capital. Cambridge tops UK for science investment as US capital surges into tech and life sciences. Report, 10 2025. URL: <https://www.enterprise.cam.ac.uk/news/cambridge-tops-uk-for-science-investment-as-us-capital-surges-into-tech-and-life-sciences/>.
- [7] E. Broda. Agentic mesh: The future of generative AI-enabled autonomous agent ecosystems. Medium, 11 2024. See also Broda, E. & Broda, D. (forthcoming). *Agentic Mesh*. O’Reilly Media. URL: <https://medium.com/@erbroda/agentic-mesh-the-future-of-generative-ai-enabled-autonomous-agent-ecosystems-d6a11381c979>.
- [8] R. S. Burt. *Structural Holes: The Social Structure of Competition*. Harvard University Press, Cambridge, MA, 1992.
- [9] Cambridge Ahead. Foreign ownership and ownership changes in the combined authority. Technical report, 2025. URL: <https://www.cambridgeahead.co.uk/wp-content/uploads/2025/04/foreign-ownership-and-ownership-changes-in-the-combined-authority.pdf>.
- [10] Cambridge Centre for Alternative Finance. New platform for interactive fintech data and visualisations, 12 2021. Describes the Cambridge FinTech Ecosystem Atlas as covering over 2,800 entities across 14 market segments. URL: <https://www.jbs.cam.ac.uk/2021/new-platform-for-interactive-fintech-data-and-visualisations/>.

- [11] Cambridge Centre for Alternative Finance. Cambridge fintech ecosystem atlas: Methodology, 2026. URL: <https://ccaf.io/atlas/methodology>.
- [12] Cambridge Enterprise. Innovation. partnership. impact. Cambridge Enterprise annual review 2023–24. Technical report, University of Cambridge, 2024. URL: <https://www.enterprise.cam.ac.uk/wp-content/uploads/2025/03/Cambridge-Enterprise-Annual-Review-2023-24.pdf>.
- [13] Cambridge Innovation Capital. Fund III and opportunity fund documentation. Technical report, 2025. See also British Business Bank case study on CIC. URL: <https://www.cic.vc/>.
- [14] Cambridge Judge Business School. Cambridge fintech and regulatory innovation accelerator, 2026. URL: <https://www.jbs.cam.ac.uk/faculty-research/centres/alternative-finance/courses/fintech-regulatory-innovation/>.
- [15] CB Insights. Insurtech 50. Technical report, 10 2025. URL: <https://www.cbinsights.com/research/report/insurtech-50-2025/>.
- [16] CB Insights. State of insurtech 2025 report. Technical report, 2 2026. Includes 2026 outlook and predictions. URL: <https://www.cbinsights.com/research/report/insurtech-trends-2025/>.
- [17] S. Chesnokova. Cambridge cracks the code: \$2.3B VC boom, deep tech surge — but diversity gaps remain. Report, 2025. URL: <https://techfundingnews.com/cambridge-cracks-the-code-2-3b-vc-boom-deep-tech-surge-but-diversity-gaps-remain/>.
- [18] T. Y. Choi, K. J. Dooley, and M. Rungtusanatham. Supply networks and complex adaptive systems: Control versus emergence. *Journal of Operations Management*, 19(3):351–366, 2001. doi:10.1016/S0272-6963(00)00068-1.
- [19] P. Cooke. Regional innovation systems: Competitive regulation in the new Europe. *Geoforum*, 23(3):365–382, 1992. doi:10.1016/0016-7185(92)90048-9.
- [20] Dealroom.co and Founders at the University of Cambridge. Cambridge tech ecosystem report 2025, 2025. Dealroom analysis valuing Cambridge’s tech ecosystem at around \$191bn and profiling it as a science-led deep tech and life sciences hub, without a separate fintech sector breakdown. URL: <https://dealroom.co/uploaded/2025/05/Cambridge-report-2025-2.pdf>.
- [21] Deloitte. The agentic reality check: Preparing for a silicon-based workforce, 12 2025. Chapter in *Tech Trends 2026*. Cites Gartner: “over 40% of agentic AI projects will fail by 2027 because legacy systems can’t support modern AI execution demands.” Identifies three infrastructure obstacles: legacy integration, data architecture, governance. URL: <https://www.deloitte.com/us/en/insights/topics/technology-management/tech-trends/2026/agentic-ai-strategy.html>.

- [22] Department for Science, Innovation and Technology. Artificial intelligence sector study 2024. Technical report, Department for Science, Innovation and Technology, Sept. 2025. URL: <https://www.gov.uk/government/publications/artificial-intelligence-sector-study-2024/artificial-intelligence-sector-study-2024>.
- [23] Department for Science, Innovation and Technology. Cambridge supercomputer set to get 6 times more powerful as government backs british ai innovation, Jan. 2026. URL: <https://www.gov.uk/government/news/cambridge-supercomputer-set-to-get-6-times-more-powerful-as-government-backs-british-ai-innovation>.
- [24] Digital Strategy Institute. Latest trends in digital business models and ecosystem innovation, 2025. URL: <https://www.digitalstrategyinstitute.org/post/latest-trends-in-digital-business-models-and-ecosystem-innovation-2025>.
- [25] Financial Conduct Authority. Mutuals public register: Society 31786, 2026. Official FCA Mutuals Public Register entry for The Venture Society Limited. URL: <https://mutuals.fca.org.uk/Search/Society/31786>.
- [26] Gartner. Gartner predicts 40% of enterprise apps will feature task-specific AI agents by 2026, up from less than 5% in 2025. Technical report, Gartner Inc., 8 2025. URL: <https://www.gartner.com/en/newsroom/press-releases/2025-08-26-gartner-predicts-40-percent-of-enterprise-apps-will-feature-task-specific-ai-agents-by-2026-up-from-less-than-5-percent-in-2025>.
- [27] A. Gawer and M. A. Cusumano. Industry platforms and ecosystem innovation. *Journal of Product Innovation Management*, 31(3):417–433, 2014. doi:10.1111/jpim.12105.
- [28] Global Cambridge Institute. GCI strategy plan. Technical report, Global Cambridge Institute, 2025. URL: <https://www.globalcambridge.org/strategy/>.
- [29] Grand View Research. AI agents market size, share & trends analysis report 2025–2033. Technical report, Grand View Research, 2025. URL: <https://www.grandviewresearch.com/industry-analysis/ai-agents-market-report>.
- [30] O. Granstrand and M. Holgersson. Innovation ecosystems: A conceptual review and a new definition. *Technovation*, 90–91:102098, 2020. doi:10.1016/j.technovation.2019.102098.
- [31] J. Haskel and S. Westlake. The UK innovation economy: A hidden gem. Technical report, 2025. Reported via Orbit RRI; underinvestment estimated at up to £143bn per year. See also Haskel & Westlake, *Capitalism without Capital* (2017) and *Restarting the Future* (2022). URL: <https://orbit-rri.org/general-news-updates/uk-innovation-under-valued-global-asset/>.
- [32] House of Commons Library. Service industries: Economic indicator, 2025. URL: <https://researchbriefings.files.parliament.uk/documents/SN02786/SN02786.pdf>.

- [33] M. Hülsmann and K. Windt, editors. *Understanding Autonomous Cooperation and Control in Logistics: The Impact of Autonomy on Management, Information, Communication and Material Flow*. Springer, Berlin, 2007. doi:10.1007/978-3-540-47450-0.
- [34] Innovate Cambridge, Beahurst, Cambridge Enterprise, and Cambridge Innovation Capital. Cambridge tops UK for science investment as US capital surges into tech and life sciences, 10 2025. International investors involved in nearly 40% of Cambridge deals; ecosystem grown from 473 to 848 active companies 2015–2025. URL: <https://innovatecambridge.com/cambridge-tops-uk-for-science-investment-as-us-capital-surges-into-tech-and-life-sciences/>.
- [35] Innovate Cambridge and Cambridge Enterprise. Annual summit data and reports. Technical report, 2025. URL: <https://innovatecambridge.com/>.
- [36] Investopedia Staff. British chip designer Arm debuts on Nasdaq in biggest US IPO since 2021. *Investopedia*, 09 2023. Accessed 12 April 2026. URL: <https://www.investopedia.com/british-chip-designer-arm-debuts-on-nasdaq-in-biggest-us-ipo-since-2021-7969648>.
- [37] D. J. Isenberg. The big idea: How to start an entrepreneurial revolution. *Harvard Business Review*, 6 2010. URL: <https://hbr.org/2010/06/the-big-idea-how-to-start-an-entrepreneurial-revolution>.
- [38] S. A. Kauffman. *The Origins of Order: Self-Organization and Selection in Evolution*. Oxford University Press, 1993.
- [39] MIT Sloan Management Review. Platform strategy summit, 2025. URL: <https://sloanreview.mit.edu/>.
- [40] OECD. Artificial intelligence and competitive dynamics in downstream markets. Technical Report 331, OECD Publishing, 2025. URL: https://www.oecd.org/en/publications/artificial-intelligence-and-competitive-dynamics-in-downstream-markets_ccf0624a-en.html.
- [41] OECD and IDIA. Innovation ecosystem strengthening learning journey. Technical report, OECD, 10 2025. URL: <https://www.oecd.org/>.
- [42] E. Ostrom. *Governing the Commons: The Evolution of Institutions for Collective Action*. Cambridge University Press, Cambridge, UK, 1990. doi:10.1017/CB09780511807763.
- [43] G. G. Parker, M. W. Van Alstyne, and S. P. Choudary. *Platform Revolution: How Networked Markets Are Transforming the Economy—and How to Make Them Work for You*. W.W. Norton, 2016.
- [44] Y. Proselkov and A. Abubakar. Systematising startup creation. Technical report, Global Cambridge Institute, 2025. December 2025; revised March 2026. URL: <https://www.globalcambridge.org/reports/>.

- [45] PwC. Insurance in 2025 and beyond. Technical report, PwC, 2022. Estimates global protection gap at US\$1.86 trillion by 2025. URL: <https://www.pwc.com/bm/en/press-releases/insurance-in-2025-and-beyond.html>.
- [46] PwC. PwC’s 2025 Responsible AI survey: From policy to practice, 10 2025. Survey of 310 US business leaders. About 61% report being at the strategic or embedded stage of Responsible AI maturity; 21% at the training stage; 18% in early stages. URL: <https://www.pwc.com/us/en/tech-effect/ai-analytics/responsible-ai-survey.html>.
- [47] PwC. PwC’s AI agent survey. Technical report, PwC, 5 2025. Survey of 300 senior US executives, May 2025. 79% report AI agents being adopted in their companies; 88% plan to increase AI-related budgets in next 12 months. URL: <https://www.pwc.com/us/en/tech-effect/ai-analytics/ai-agent-survey.html>.
- [48] G. Sarafin. Seven business models for creating ecosystem value, 2021. Updated and republished 2024. URL: https://www.ey.com/en_gl/alliances/seven-business-models-for-creating-ecosystem-value.
- [49] A. Shipilov and F. Burelli. Don’t confuse platforms with ecosystems, 2022. URL: <https://knowledge.insead.edu/strategy/dont-confuse-platforms-ecosystems>.
- [50] D. Sjödin, V. Parida, M. Palmié, and J. Wincent. How AI capabilities enable business model innovation: Scaling AI through co-evolutionary processes and feedback loops. *Journal of Business Research*, 134:574–587, 2021. doi:10.1016/j.jbusres.2021.05.009.
- [51] J. S. Srari, M. Gregory, et al. Capturing value from global networks. Technical report, Centre for International Manufacturing, Institute for Manufacturing, University of Cambridge, 2014. URL: https://www.ifm.eng.cam.ac.uk/uploads/Resources/Reports/Capturing_Value_from_Global_Networks_web.pdf.
- [52] E. Stam and B. Spigel. Entrepreneurial ecosystems and regional policy. In R. Blackburn, D. D. Clercq, J. Heinonen, and Z. Wang, editors, *The SAGE Handbook of Small Business and Entrepreneurship*. SAGE, London, 2018. Working paper version dated 2016 (Utrecht School of Economics WP 16-13); chapter published in 2018. doi:10.4135/9781473984080.n21.
- [53] E. Stam and A. van de Ven. Entrepreneurial ecosystem elements. *Small Business Economics*, 56(2):809–832, 2021. doi:10.1007/s11187-019-00270-6.
- [54] A. Sukharevsky, A. Krivkovich, A. Gast, A. Storozhev, D. Maor, D. Mahadevan, L. Hämäläinen, and S. Durth. The agentic organization: Contours of the next paradigm for the AI era. Technical report, McKinsey & Company, 9 2025. URL: <https://www.mckinsey.com/capabilities/people-and-organizational-performance/our-insights/the-agentic-organization-contours-of-the-next-paradigm-for-the-ai-era>.
- [55] University of Cambridge. 10 cambridge ai spinouts, Apr. 2025. URL: <https://www.cam.ac.uk/stories/cambridge-ai-spinouts>.

- [56] University of Cambridge. Annual report and financial statements 2024–25. Technical report, University of Cambridge, 2025. URL: https://www.cam.ac.uk/system/files/university_of_cambridge_group_annual_reports_financial_statements_2024-25.pdf.
- [57] F. Wan et al. Scaling high and wide: Leveraging AI for ecosystem scale-up. In *Academy of Management Proceedings*, 2025. Conference paper; publication details not independently verified.
- [58] World Intellectual Property Organization. Global innovation index 2024: Unlocking the promise of social entrepreneurship. Technical report, Geneva, 2024. UK ranks 5th overall on the GII 2024. URL: https://www.wipo.int/web-publications/global-innovation-index-2024/assets/67729/2000%20Global%20Innovation%20Index%202024_WEB31ite.pdf, doi:10.34667/tind.50062.
- [59] L. Xu, S. Mak, M. Minaricova, and A. Brintrup. On implementing autonomous supply chains: A multi-agent system approach. *Computers in Industry*, 161:104120, 2024. doi:10.1016/j.compind.2024.104120.
- [60] L. Xu, S. Mak, Y. Proselkov, and A. Brintrup. Towards autonomous supply chains: Definition, characteristics, conceptual framework, and autonomy levels. *Journal of Industrial Information Integration*, 42:100698, 2024. doi:10.1016/j.jii.2024.100698.